

Introduction to PyTorch

This tutorial shows the basics of PyTorch library.

We design simple Neural Networks for classification task on MNIST dataset.

You're probably going to need PyTorch documentation:

<https://pytorch.org/docs/stable/index.html>

and tutorials:

https://pytorch.org/tutorials/beginner/basics/quickstart_tutorial.html

```
In [ ]: !mkdir data
```

1. PyTorch is based on tensor operations. First, let's try to use them:

- create simple python list with four values and convert it to PyTorch tensor
- create numpy array with random values and shape (1,3,7,7) and convert it to Pytorch tensor
- create PyTorch tensor with random values and shape (1,3,7,7) with preset seed
- create PyTorch tensor with linear space in range from -5 to 15 and reshape it to tensor with shape (1,3,7,7)
- create PyTorch tensor of zeros with shape (1,3,7,7)

For display use `print` function.

```
In [17]: import torch
import numpy as np
from typing import Tuple
import tqdm
from tabulate import tabulate
python_list = [1, 2, 3, 4]
tensor_from_list = torch.tensor(python_list)
print(f"Tensor from list: {tensor_from_list}")

numpy_arr = np.random.rand(1, 3, 7, 7)
tensor_from_numpy = torch.tensor(numpy_arr)
print(f"Tensor from numpy arr: {tensor_from_numpy}")

# dzięki temu przy kazdym uruchomieniu otrzymujemy te same losowe wartości
random_tensor = torch.rand(1, 3, 7, 7)
torch.manual_seed(0)
print(f"Random tensor with preset seed: {random_tensor}")
```

```
# linspace tworzy nam liniowe rozłoenie danych (czyli chyba ciąg? taka sa
linear_tensor = torch.linspace(-5, 15, steps=1*3*7*7).reshape(1, 3, 7, 7)
print(linear_tensor)

zeros_tensor = torch.zeros(1, 3, 7, 7)
print(f"Zeros tensor: {zeros_tensor}")
```

```
Tensor from list: tensor([1, 2, 3, 4])
Tensor from numpy arr: tensor([[[[0.1979, 0.2089, 0.8812, 0.3688, 0.5101,
0.0525, 0.4303],
[0.9911, 0.0207, 0.5338, 0.7258, 0.9917, 0.8546, 0.1344],
[0.4611, 0.9622, 0.4394, 0.2708, 0.8185, 0.9220, 0.1635],
[0.6040, 0.1079, 0.8381, 0.2752, 0.3549, 0.1486, 0.9634],
[0.0252, 0.1644, 0.1698, 0.5441, 0.2444, 0.1428, 0.9955],
[0.5414, 0.9531, 0.8055, 0.7594, 0.8380, 0.9566, 0.1384],
[0.4404, 0.6970, 0.7797, 0.6291, 0.5225, 0.2466, 0.7973]],
[[0.1731, 0.0968, 0.3959, 0.6689, 0.0241, 0.9399, 0.8844],
[0.4737, 0.0414, 0.4566, 0.8066, 0.3761, 0.5983, 0.3666],
[0.5807, 0.8480, 0.1096, 0.3851, 0.8597, 0.7061, 0.1032],
[0.8467, 0.8615, 0.8638, 0.6549, 0.5167, 0.6238, 0.3705],
[0.5182, 0.1182, 0.4048, 0.0111, 0.2754, 0.1509, 0.0785],
[0.5646, 0.5618, 0.0799, 0.8472, 0.8155, 0.0692, 0.0155],
[0.1041, 0.2831, 0.9905, 0.5570, 0.4696, 0.5949, 0.6720]],
[[0.2689, 0.1559, 0.3610, 0.9099, 0.0967, 0.6367, 0.9011],
[0.8516, 0.7169, 0.7008, 0.7046, 0.2938, 0.7228, 0.4940],
[0.0229, 0.3864, 0.5015, 0.7872, 0.9237, 0.8596, 0.1773],
[0.4732, 0.1491, 0.9773, 0.0663, 0.6229, 0.2485, 0.6849],
[0.9634, 0.5058, 0.4425, 0.5908, 0.0383, 0.6257, 0.2074],
[0.2559, 0.3446, 0.8127, 0.1353, 0.5073, 0.6480, 0.5414],
[0.4851, 0.0487, 0.6537, 0.9839, 0.3720, 0.2927, 0.2682]]]],
dtype=torch.float64)
Random tensor with preset seed: tensor([[[[0.4963, 0.7682, 0.0885, 0.1320,
0.3074, 0.6341, 0.4901],
[0.8964, 0.4556, 0.6323, 0.3489, 0.4017, 0.0223, 0.1689],
[0.2939, 0.5185, 0.6977, 0.8000, 0.1610, 0.2823, 0.6816],
[0.9152, 0.3971, 0.8742, 0.4194, 0.5529, 0.9527, 0.0362],
[0.1852, 0.3734, 0.3051, 0.9320, 0.1759, 0.2698, 0.1507],
[0.0317, 0.2081, 0.9298, 0.7231, 0.7423, 0.5263, 0.2437],
[0.5846, 0.0332, 0.1387, 0.2422, 0.8155, 0.7932, 0.2783]],
[[0.4820, 0.8198, 0.9971, 0.6984, 0.5675, 0.8352, 0.2056],
[0.5932, 0.1123, 0.1535, 0.2417, 0.7262, 0.7011, 0.2038],
[0.6511, 0.7745, 0.4369, 0.5191, 0.6159, 0.8102, 0.9801],
[0.1147, 0.3168, 0.6965, 0.9143, 0.9351, 0.9412, 0.5995],
[0.0652, 0.5460, 0.1872, 0.0340, 0.9442, 0.8802, 0.0012],
[0.5936, 0.4158, 0.4177, 0.2711, 0.6923, 0.2038, 0.6833],
[0.7529, 0.8579, 0.6870, 0.0051, 0.1757, 0.7497, 0.6047]],
[[0.1100, 0.2121, 0.9704, 0.8369, 0.2820, 0.3742, 0.0237],
[0.4910, 0.1235, 0.1143, 0.4725, 0.5751, 0.2952, 0.7967],
```

```

[0.1957, 0.9537, 0.8426, 0.0784, 0.3756, 0.5226, 0.5730],
[0.6186, 0.6962, 0.5300, 0.2560, 0.7366, 0.0204, 0.2036],
[0.3748, 0.2564, 0.3251, 0.0902, 0.3936, 0.6069, 0.1743],
[0.4743, 0.8579, 0.4486, 0.5139, 0.4569, 0.6012, 0.8179],
[0.9736, 0.8175, 0.9747, 0.4638, 0.0508, 0.2630, 0.8405]]]])
tensor([[[[-5.0000, -4.8630, -4.7260, -4.5890, -4.4521, -4.3151, -4.1781],
[-4.0411, -3.9041, -3.7671, -3.6301, -3.4932, -3.3562, -3.2192],
[-3.0822, -2.9452, -2.8082, -2.6712, -2.5342, -2.3973, -2.2603],
[-2.1233, -1.9863, -1.8493, -1.7123, -1.5753, -1.4384, -1.3014],
[-1.1644, -1.0274, -0.8904, -0.7534, -0.6164, -0.4795, -0.3425],
[-0.2055, -0.0685, 0.0685, 0.2055, 0.3425, 0.4795, 0.6164],
[ 0.7534, 0.8904, 1.0274, 1.1644, 1.3014, 1.4384, 1.575

3]],

[[ 1.7123, 1.8493, 1.9863, 2.1233, 2.2603, 2.3973, 2.5342],
[ 2.6712, 2.8082, 2.9452, 3.0822, 3.2192, 3.3562, 3.4932],
[ 3.6301, 3.7671, 3.9041, 4.0411, 4.1781, 4.3151, 4.4521],
[ 4.5890, 4.7260, 4.8630, 5.0000, 5.1370, 5.2740, 5.4110],
[ 5.5479, 5.6849, 5.8219, 5.9589, 6.0959, 6.2329, 6.3699],
[ 6.5068, 6.6438, 6.7808, 6.9178, 7.0548, 7.1918, 7.3288],
[ 7.4658, 7.6027, 7.7397, 7.8767, 8.0137, 8.1507, 8.287

7]],

[[ 8.4247, 8.5616, 8.6986, 8.8356, 8.9726, 9.1096, 9.2466],
[ 9.3836, 9.5205, 9.6575, 9.7945, 9.9315, 10.0685, 10.2055],
[10.3425, 10.4795, 10.6164, 10.7534, 10.8904, 11.0274, 11.1644],
[11.3014, 11.4384, 11.5753, 11.7123, 11.8493, 11.9863, 12.1233],
[12.2603, 12.3973, 12.5342, 12.6712, 12.8082, 12.9452, 13.0822],
[13.2192, 13.3562, 13.4932, 13.6301, 13.7671, 13.9041, 14.0411],
[14.1781, 14.3151, 14.4521, 14.5890, 14.7260, 14.8630, 15.000

0]]]])
Zeros tensor: tensor([[[[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.]],

[[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.]],

[[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0.]]],

```

2. PyTorch allows applying GPU for computations. Check if gpu (CUDA) is available and then use it as `device` , else use `'cpu'` . Then, move one of your tensors to selected device.

```
# przenosimy nasz tensor na urządzenie które będzie korzystać z wcześniej
random_tensor = random_tensor.to(device)
print(f"Tensor on {device}: {random_tensor}")
```

Using device: cpu

```
Tensor on cpu: tensor([[[[7.0080e-01, 8.6437e-01, 7.5487e-01, 7.1401e-01,
4.8101e-01,
    4.3614e-01, 6.4462e-01],
[9.5264e-01, 1.2725e-01, 2.3432e-01, 7.9959e-01, 5.4008e-01,
    8.4314e-02, 2.6639e-01],
[7.7565e-02, 1.8192e-01, 6.4973e-01, 5.9840e-01, 5.7920e-01,
    5.0073e-02, 3.5495e-01],
[4.1902e-05, 5.5996e-01, 9.1004e-01, 2.8843e-01, 2.7448e-01,
    1.8270e-01, 5.3634e-01],
[2.0827e-01, 9.1346e-01, 3.7689e-01, 7.4214e-01, 3.2981e-01,
    2.3370e-01, 7.7565e-01],
[7.7355e-01, 9.7310e-01, 6.6742e-01, 5.7325e-01, 6.4083e-01,
    7.8802e-01, 2.8502e-01],
[4.1989e-01, 8.2637e-01, 3.2067e-01, 1.5849e-01, 6.3668e-01,
    8.8936e-01, 7.9497e-01]]],
[[[6.8548e-01, 7.9268e-01, 5.1928e-01, 7.5358e-01, 9.9550e-01,
    8.8905e-01, 4.2463e-01],
[1.7213e-01, 8.9372e-01, 1.1398e-01, 1.0622e-01, 8.0143e-01,
    8.8916e-01, 9.6725e-02],
[4.1713e-01, 4.8904e-01, 1.7239e-02, 2.6127e-01, 7.6730e-01,
    4.9176e-01, 3.8155e-01],
[2.3672e-02, 2.2816e-01, 2.9710e-01, 7.7397e-01, 2.7208e-01,
    8.5644e-01, 5.5000e-01],
[3.6513e-01, 3.6393e-01, 1.2503e-01, 9.5675e-01, 1.7554e-01,
    6.4664e-01, 5.2367e-01],
[4.8684e-01, 5.2972e-01, 5.4134e-01, 4.2196e-01, 7.8296e-01,
    2.1832e-01, 8.7881e-01],
[2.8872e-01, 1.6445e-01, 4.0563e-01, 4.6850e-01, 2.4356e-01,
    4.7162e-01, 4.3645e-01]]],
[[[9.6860e-01, 3.9752e-01, 2.4091e-01, 3.7410e-01, 2.2564e-01,
    5.2210e-01, 8.2195e-01],
[8.2360e-01, 6.2337e-01, 1.4158e-01, 3.2934e-01, 1.1538e-01,
    1.1906e-01, 7.0142e-01],
[3.6576e-01, 4.5279e-01, 9.2385e-02, 5.0944e-01, 5.9279e-01,
    9.1178e-01, 9.3597e-01],
[2.7988e-01, 2.5921e-02, 7.8537e-01, 9.3055e-01, 8.2178e-01,
    9.0187e-01, 3.3890e-01],
[1.6154e-01, 2.4392e-01, 4.9567e-01, 6.8660e-01, 8.9719e-02,
    6.1184e-01, 8.4940e-01],
[9.7104e-01, 8.9192e-01, 1.7125e-01, 4.7201e-02, 7.9129e-01,
    3.6652e-01, 6.1345e-01],
[1.6851e-01, 6.0670e-02, 3.3917e-01, 6.6320e-01, 5.6449e-01,
    2.7340e-01, 3.8840e-01]]]])
```

3. To train a network, we need a dataset.

Please download `MNIST` dataset with `torchvision.dataset`.

For any kind of ML task, validation or testing is required.

So, create train and test datasets.

For train dataset apply also augmentation transforms, crop, translation and rotation.

For both apply ToTensor.

Next, pack datasets into `DataLoader` s with batch size of 64. Use variables with names: `train_loader` and `test_loader` .

Next display sizes of datasets, shapes of elements and display a few images and their labels.

Finally, compare the number of objects in each class in both datasets.

```
In [3]: import torch
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor, Compose, RandomCrop, RandomAffine
import matplotlib.pyplot as plt

# Compose łączy kilka transformacji w jedną sekwencję
# RandomCrop przycina obraz do rozmiaru 28x28 pikseli z losowym paddingiem
# RandomAffine losowa rotacja obrazu o maks 15 stopni i translacja czyli
# ToTensor konwertuje na tensor bo tak chce torch
train_transform = Compose([
    RandomCrop(28, padding=4),
    RandomAffine(degrees=15, translate=(0.1, 0.1)),
    ToTensor()
])

# Tu nie stosujemy augmentacji zeby wyniki byly porównywalne
test_transform = Compose([
    ToTensor()
])

try:
    train_dataset = datasets.MNIST(
        root="./data", train=True, download=True, transform=train_transform
    )
    test_dataset = datasets.MNIST(
        root="./data", train=False, download=True, transform=test_transform
    )
except Exception as e:
    print(f"Failed to download MNIST dataset: {e}")

# DataLoader klasa ułatwia iteracji po danych w partiach
# batch_size liczba probek w jednej partii
# shuffle losowe mieszanie danych w kazdej epoce
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

print(f"Train dataset size: {len(train_dataset)}")
```

```
print(f"Test dataset size: {len(test_dataset)}")
```

Train dataset size: 60000

Test dataset size: 10000

```
In [4]: dataiter = iter(train_loader)
images, labels = next(dataiter)
print(f"Batch shape: {images.size()}") # [64, 1, 28, 28] - 64 liczba obra
print(f"Labels shape: {labels.size()}") # 64 etykiety

fig, axes = plt.subplots(1, 6, figsize=(15, 4))
for i in range(6):
    ax = axes[i]
    ax.imshow(images[i].numpy().squeeze(), cmap="gray")
    ax.set_title(f"Label: {labels[i].item()}")
    ax.axis("off")
plt.show()


train_labels = [label for _, label in train_loader.dataset]
test_labels = [label for _, label in test_loader.dataset]

# zlicza liczbę wystąpień każdej klasy
train_class_counts = torch.tensor(train_labels).bincount()
test_class_counts = torch.tensor(test_labels).bincount()

print(f"Train class counts: {train_class_counts}")
print(f"Test class counts: {test_class_counts}")
```

Batch shape: torch.Size([64, 1, 28, 28])

Labels shape: torch.Size([64])



Train class counts: tensor([5923, 6742, 5958, 6131, 5842, 5421, 5918, 6265, 5851, 5949])

Test class counts: tensor([980, 1135, 1032, 1010, 982, 892, 958, 1028, 974, 1009])

```
In [20]: import matplotlib.pyplot as plt
```

```
def draw_loss_test(epochs, history):
    loss_train = history["loss_train"]
    loss_test = history["loss_test"]
    acc_train = history["acc_train"]
    acc_test = history["acc_test"]

    loss_train_shape = len(loss_train)
    loss_test_shape = len(loss_test)
    acc_train_shape = len(acc_train)
    acc_test_shape = len(acc_test)
```

```

if (
    loss_train_shape != loss_test_shape
    or acc_train_shape != acc_test_shape
):
    raise ValueError(
        f"Different number of epochs for train and test loss: {loss_t

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(epochs, history["loss_train"], label="Train loss")
plt.plot(epochs, history["loss_test"], label="Test loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, history["acc_train"], label="Train accuracy")
plt.plot(epochs, history["acc_test"], label="Test accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()

plt.show()

```

5. We have our dataset ready, let's create model for classification task.

Please define class `MLP` as Multi Layer Perceptron with two hidden fully connected layers with bias.

Class must inherit from `torch.nn.Module`.

Apply following configuration:

- first layer with 512 neurons,
- second layer with 512 neurons,
- output layer adjusted to the size of classification problem.

For `__init__` method add parameters: `input_shape` and `output_size`.

Don't forget about nonlinearities!

For hidden layers you can use `ReLU` module from `torch.nn`.

For output apply softmax function.

Define layer-by-layer network processing in `forward` method with argument as a network.

Help: input tensor - batch of images with shape (batch_size, channels, height, width)
- (channels = 1 for gray scale images).

Instantiate model as `net` object.

Layers:

- `torch.nn.Sequential`
 - layer allows for forward pass through component layers:
`t_in: Tensor`
`t_out: Tensor = t_in`
`for L in layers:`
 `t_out = L(t_out)`
`return t_out`
- `torch.nn.Flatten`
 - layer makes input tensor flattened:
 - (bs, CH, H, W) -> (bs, CH * H * W)
- `torch.nn.Linear(ch_in, ch_out, bias)`
 - 'classical' neural network layer - fully connected
 - `ch_in` is a number of input channels
 - `ch_out` is a number of output channels / number of neurons in layer
 - `bias` - whether to use bias parameter
 - for Linear layers it is recommended to use flatten layer before, when input has more than 2 dimensions
 - operation implemented by this layer is a vector / matrix multiplication
 - `y = W x v` or `y = W x v + b`
 - `W` has a shape [`ch_out`, `ch_in`]
 - `v` has a shape [`ch_in`]
 - `b` has a shape [`ch_in`]
 - `y` has a shape [`ch_out`]
- `torch.nn.ReLU`
 - layer applies ReLU function on input tensor
- `torch.nn.Softmax(dim)`
 - layer applies softmax function on input tensor
 - `dim` - dimension over which function is calculated

For the formulas of activation function go to [torch documentation](#).

```
In [6]: import torch.nn as nn
```

```
# klasa definiująca wielowarstwowy perceptron
class MLP(nn.Module):

    def __init__(self, input_shape, output_size) -> None:
        super(MLP, self).__init__()
        # Sequential to kontener który pozwala na zdefiniowanie listy war
        self.model = nn.Sequential(
            # Flatten warstwa spłaszczająca tensor z (batch_size, channel
            # np. obrazki 28x28 maja (batch_size, 1, 28, 28) i spłaszczy
            # 784 to liczba cech
            nn.Flatten(),
            # Linear  $y = Wx + b$  (W macierz wag, x dane wejściowe, b wekto
            # in_features - liczba wejściowych cech, out_features - liczb
            nn.Linear(input_shape[0] * input_shape[1] * input_shape[2], 5
            # ReLU  $f(x) = \max(0, x)$ 
            # warstwa aktywacji - wprowadza nieliniowość modelu (bez tego
            # pozwala sieci na uczenie się bardziej złożonych funkcji
            nn.ReLU(),
            nn.Linear(512, 512, bias=True),
            nn.ReLU(),
            # ostatnia warstwa przekształca dane wyjściowe na dane o rozm
            nn.Linear(512, output_size, bias=True),
            # Softmax  $\text{softmax}(x_i) = \exp(x_i) / \sum(\exp(x_j))$ 
            # warstwa aktywacji - przekształca wyjścia w prawdopodobieństw
            # dim=1 oznacza że softmax jest stosowany wzdłuż osi klas (d
            # często nie używamy tego gdy korzystamy z CrossEntropyLoss bo
            nn.Softmax(dim=1)
        )

    # definiuje przepływ danych przez sieć, x: tensor wejściowy (batch_siz
    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.model(x)

# obrazek w skali szarości o rozmiarze 28x28 pikseli
input_shape = (1, 28, 28)
# liczba klas do klasyfikacji np. 10 do liczb od 0 do 9
output_size = 10

net = MLP(input_shape, output_size)
net = net.to(device)

net2 = MLP(input_shape, output_size)
net2 = net2.to(device)
```

6. To train network we need to know 'how good or bad' results it gives. Please, instantiate `torch.nn.CrossEntropyLoss` as `loss_fn`.

```
In [7]: # informuje jak dobrze model przewiduje prawdziwe klasy
loss_fn = nn.CrossEntropyLoss()
```

7. To score network define accuracy metric. For network output you need to decide

what is the final network answer. For classification we can assume, that the final answer is the class with highest probability (`argmax`).

`torch.no_grad()` prevents gradient requirement for computations inside method.

```
In [8]: from abc import ABC, abstractmethod
        from typing import Any

        class BaseMetric(ABC):

            @abstractmethod
            def __call__(self, y_pred, y_ref) -> Any:
                raise NotImplementedError()

        class AccuracyMetric(BaseMetric):

            def __init__(self) -> None:
                pass

            @torch.no_grad()
            def __call__(self, y_pred: torch.Tensor, y_ref: torch.Tensor) -> torch.Tensor:
                """
                :param y_pred: tensor of shape (batch_size, num_of_classes) type Long
                :param y_ref: tensor with shape (batch_size,) and type Long
                :return: scalar tensor with accuracy metric for batch
                """

                predicted_classes = torch.argmax(y_pred, dim=1) # zwroci indeks k
                correct_predictions = (predicted_classes == y_ref).float() # porównanie
                score: torch.Tensor = correct_predictions.mean() # srednia z poprawnych

                return score

        metric = AccuracyMetric()
```

8. To change network parameters, we need an optimizer object. Instantiate `torch.optim.SGD` (with network parameters) as `optimizer`. Use learning rate = 0.001.

```
In [9]: # Stochastic Gradient Descent algorytm do optymalizacji (minimizacji funkcji
        # straty)
        optimizer = torch.optim.SGD(net.parameters(), lr=0.01)
        optimizer2 = torch.optim.SGD(net2.parameters(), lr=0.01)
```

9. Now define training / testing function:

```
In [24]: def train_or_test(
```

```
model, # sieć neuronowa która trenujemy lub testujemy
data_generator, # generator danych, który dostarcza dane wejściowe i
criterion, # funkcja straty, która oblicza różnice między predykcjami
metric: BaseMetric, # funkcja metryki, która oblicza dokładność pred
mode: str = "test",
optimizer: torch.optim.Optimizer = None, # optymalizator – aktualizacja
update_period: int = None, # Okres aktualizacji, określa co ile batchy
device=torch.device("cpu"),
) -> Tuple[torch.nn.Module, float, float]:

    # change model mode to train or test
    if mode == 'train':
        model.train()
    elif mode == 'test':
        model.eval()
    else:
        raise RuntimeError(f"Unsupported mode: {mode}. Use 'train' or 'test'")

    # move model to device
    model.to(device)

    # reset model parameters' gradients with optimizer
    if mode == 'train':
        optimizer.zero_grad()

    total_loss = 0.0 # suma strat dla wszystkich batchy
    total_accuracy = 0.0 # suma dokładności dla wszystkich batchy
    samples_num = 0 # liczba próbek przetworzonych przez model

    for i, (X, y) in enumerate(tqdm.tqdm(data_generator)):
        # convert tensors to device
        X, y = X.to(device), y.to(device)

        # process by network
        y_pred = model(X)

        # calculate loss
        loss = criterion(y_pred, y)

        # designate gradient based on loss
        if mode == 'train':
            loss.backward()

        if mode == 'train' and (i+1) % update_period == 0:
            # update parameters with optimizer
            optimizer.step()
            optimizer.zero_grad()

        # calculate accuracy
        accuracy = metric(y_pred, y)

        total_loss += loss.item() * y_pred.shape[0] # y_pred.shape[0] to
        total_accuracy += accuracy.item() * y_pred.shape[0]
```

```

        samples_num += y_pred.shape[0]

    if samples_num == 0:
        return model, 0.0, 0.0

    return model, total_loss / samples_num, total_accuracy / samples_num

```

```

In [ ]: net_preview, loss_preview, acc_preview = train_or_test(
        net,
        train_loader,
        loss_fn, metric,
        'train',
        update_period=5,
        optimizer=optimizer,
        device=device,
    )

```

10. Prepare training loop function (over epochs):

- adjust max number of epochs to achieve satisfactory results,
- **EXTENSION EXERCISE** implement stopping training when accuracy exceeds certain value.

```

In [25]: def training(model,
                    train_loader,
                    test_loader,
                    loss_fn,
                    metric,
                    optimizer,
                    update_period,
                    epoch_max,
                    device,
                    early_stopping_accuracy=None):
    loss_train = []
    loss_test = []
    acc_train = []
    acc_test = []

    for e in range(epoch_max):
        print(f"\nEpoka: {e + 1}/{epoch_max}")

        model.train()
        total_loss_train = 0.0
        total_acc_train = 0.0
        samples_num_train = 0

        for i, (X, y) in enumerate(tqdm.tqdm(train_loader, colour="red",
                                             X, y = X.to(device), y.to(device)

                                     y_pred = model(X)

```

```
    loss = loss_fn(y_pred, y)
    loss.backward()

    if (i + 1) % update_period == 0:
        optimizer.step()
        optimizer.zero_grad()

    accuracy = metric(y_pred, y)
    batch_size = y.size(0)
    total_loss_train += loss.item() * batch_size
    total_acc_train += accuracy.item() * batch_size
    samples_num_train += batch_size

epoch_loss_train = total_loss_train / samples_num_train
epoch_acc_train = total_acc_train / samples_num_train
loss_train.append(epoch_loss_train)
acc_train.append(epoch_acc_train)

print(
    f"Train loss: {epoch_loss_train:.4f}\nTrain accuracy: {epoch_
)

model.eval()
total_loss_test = 0.0
total_acc_test = 0.0
samples_num_test = 0

with torch.no_grad():
    for X, y in tqdm.tqdm(test_loader, colour='green', smoothing=
        X, y = X.to(device), y.to(device)

        y_pred = model(X)

        loss = loss_fn(y_pred, y)
        accuracy = metric(y_pred, y)

        batch_size = y.size(0)
        total_loss_test += loss.item() * batch_size
        total_acc_test += accuracy.item() * batch_size
        samples_num_test += batch_size

epoch_loss_test = total_loss_test / samples_num_test
epoch_acc_test = total_acc_test / samples_num_test
loss_test.append(epoch_loss_test)
acc_test.append(epoch_acc_test)

print(f"Test loss: {epoch_loss_test:.4f}\nTest accuracy: {epoch_a

if (early_stopping_accuracy and epoch_acc_train >= early_stopping
    print(f"Training accuracy of {epoch_acc_train:.4f} achieved,
    break

return model, {'loss_train': loss_train,
```

```
'acc_train': acc_train,  
'loss_test': loss_test,  
'acc_test': acc_test}
```

```
In [10]: def train_one_epoch(  
    model, train_loader, loss_fn, metric, optimizer, update_period, device  
    ):  
    model.train()  
    total_loss_train = 0.0  
    total_acc_train = 0.0  
    samples_num_train = 0  
  
    with tqdm.tqdm(  
        train_loader, colour="red", ncols=100  
    ) as t:  
        for i, (X, y) in enumerate(t):  
            X, y = X.to(device), y.to(device)  
  
            y_pred = model(X)  
            loss = loss_fn(y_pred, y)  
            loss.backward()  
  
            if (i + 1) % update_period == 0:  
                optimizer.step()  
                optimizer.zero_grad()  
  
                accuracy = metric(y_pred, y)  
                batch_size = y.size(0)  
                total_loss_train += loss.item() * batch_size  
                total_acc_train += accuracy.item() * batch_size  
                samples_num_train += batch_size  
  
                current_loss = total_loss_train / samples_num_train  
                current_acc = total_acc_train / samples_num_train  
  
                t.set_postfix(loss=f'{current_loss:.4f}', accuracy=f'{current  
  
epoch_loss_train = total_loss_train / samples_num_train  
epoch_acc_train = total_acc_train / samples_num_train  
  
    return epoch_loss_train, epoch_acc_train  
  
def test_one_epoch(model, test_loader, loss_fn, metric, device):  
    model.eval()  
    total_loss_test = 0.0  
    total_acc_test = 0.0  
    samples_num_test = 0  
  
    with torch.no_grad():  
        with tqdm.tqdm(  
            test_loader, colour="green", ncols=100  
        ) as t:
```

```
        for X, y in t:
            X, y = X.to(device), y.to(device)

            y_pred = model(X)
            loss = loss_fn(y_pred, y)
            accuracy = metric(y_pred, y)

            batch_size = y.size(0)
            total_loss_test += loss.item() * batch_size
            total_acc_test += accuracy.item() * batch_size
            samples_num_test += batch_size

            current_loss = total_loss_test / samples_num_test
            current_acc = total_acc_test / samples_num_test

            t.set_postfix(loss=f'{current_loss:.4f}', accuracy=f'{cur

epoch_loss_test = total_loss_test / samples_num_test
epoch_acc_test = total_acc_test / samples_num_test

return epoch_loss_test, epoch_acc_test

def test_or_train(
    model,
    train_loader,
    test_loader,
    loss_fn,
    metric,
    optimizer,
    update_period,
    epoch_max,
    device,
    mode="train",
    early_stopping_accuracy=None,
):
    loss_train = []
    loss_test = []
    acc_train = []
    acc_test = []

    for e in range(epoch_max):
        print(f"Epoch: {e + 1}/{epoch_max}")

        if mode in ["train", "both"]:
            epoch_loss_train, epoch_acc_train = train_one_epoch(
                model,
                train_loader,
                loss_fn,
                metric,
                optimizer,
                update_period,
                device,
```



```

    )
    loss_train.append(epoch_loss_train)
    acc_train.append(epoch_acc_train)

    if mode in ["test", "both"]:
        epoch_loss_test, epoch_acc_test = test_one_epoch(
            model, test_loader, loss_fn, metric, device
        )
        loss_test.append(epoch_loss_test)
        acc_test.append(epoch_acc_test)

    if (
        early_stopping_accuracy
        and epoch_acc_train >= early_stopping_accuracy
    ):
        print(
            f"Training accuracy of {epoch_acc_train:.4f} achieved, st
        )
        break

    return model, {
        "loss_train": loss_train,
        "acc_train": acc_train,
        "loss_test": loss_test,
        "acc_test": acc_test,
    }

```


11. Display training history.

```

In [13]: net2, history2 = test_or_train(
    net2,
    train_loader,
    test_loader,
    loss_fn,
    metric,
    optimizer2,
    update_period=5,
    epoch_max=100,
    device=device,
    mode="both",
    early_stopping_accuracy=0.98,
)



```



Epoch: 1/100



100%|  | 938/938 [00:04<00:00, 198.24it/s, accuracy=0.1132, loss=2.3024]



100%|  | 157/157 [00:00<00:00, 500.36it/s, accuracy=0.1276, loss=2.3016]



Epoch: 2/100



100%|  | 938/938 [00:04<00:00, 202.14it/s, accuracy=0.1279, loss=2.3018]
100%|  | 157/157 [00:00<00:00, 478.90it/s, accuracy=0.1411, loss=2.3004]
Epoch: 3/100



100%|  | 938/938 [00:04<00:00, 203.05it/s, accuracy=0.1262, loss=2.3011]
100%|  | 157/157 [00:00<00:00, 510.15it/s, accuracy=0.1157, loss=2.2990]
Epoch: 4/100


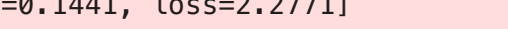
100%|  | 938/938 [00:04<00:00, 202.04it/s, accuracy=0.1097, loss=2.3002]
100%|  | 157/157 [00:00<00:00, 501.29it/s, accuracy=0.0997, loss=2.2971]
Epoch: 5/100


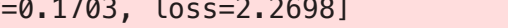
100%|  | 938/938 [00:04<00:00, 200.40it/s, accuracy=0.0999, loss=2.2990]
100%|  | 157/157 [00:00<00:00, 484.43it/s, accuracy=0.0980, loss=2.2940]
Epoch: 6/100


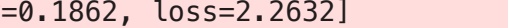
100%|  | 938/938 [00:04<00:00, 198.90it/s, accuracy=0.0987, loss=2.2967]
100%|  | 157/157 [00:00<00:00, 488.73it/s, accuracy=0.0980, loss=2.2872]
Epoch: 7/100



100%|  | 938/938 [00:04<00:00, 198.49it/s, accuracy=0.0987, loss=2.2917]
100%|  | 157/157 [00:00<00:00, 495.19it/s, accuracy=0.0980, loss=2.2727]
Epoch: 8/100



100%|  | 938/938 [00:04<00:00, 198.98it/s, accuracy=0.1048, loss=2.2844]
100%|  | 157/157 [00:00<00:00, 496.70it/s, accuracy=0.1703, loss=2.2576]
Epoch: 9/100



100%|  | 938/938 [00:04<00:00, 189.19it/s, accuracy=0.1441, loss=2.2771]
100%|  | 157/157 [00:00<00:00, 494.45it/s, accuracy=0.2138, loss=2.2443]
Epoch: 10/100



100%|  | 938/938 [00:04<00:00, 197.21it/s, accuracy=0.1703, loss=2.2698]
100%|  | 157/157 [00:00<00:00, 495.92it/s, accuracy=0.2413, loss=2.2335]
Epoch: 11/100



100%|  | 938/938 [00:04<00:00, 197.04it/s, accuracy=0.1862, loss=2.2632]
100%|  | 157/157 [00:00<00:00, 490.07it/s, accuracy=0.2526, loss=2.2246]
Epoch: 12/100



100%|  | 938/938 [00:04<00:00, 191.55it/s, accuracy=0.1989, loss=2.2553]
100%|  | 157/157 [00:00<00:00, 376.70it/s, accuracy=0.2598, loss=2.2138]
Epoch: 13/100



100%|  | 938/938 [00:04<00:00, 187.97it/s, accuracy=0.2107, loss=2.2469]
100%|  | 157/157 [00:00<00:00, 474.79it/s, accuracy=0.2620, loss=2.1975]
Epoch: 14/100


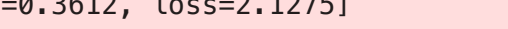
100%|  | 938/938 [00:04<00:00, 191.66it/s, accuracy=0.2213, loss=2.2339]
100%|  | 157/157 [00:00<00:00, 374.21it/s, accuracy=0.2765, loss=2.1702]
Epoch: 15/100


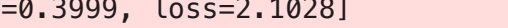
100%|  | 938/938 [00:04<00:00, 197.28it/s, accuracy=0.2364, loss=2.2171]
100%|  | 157/157 [00:00<00:00, 485.69it/s, accuracy=0.3091, loss=2.1412]
Epoch: 16/100


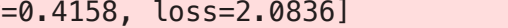
100%|  | 938/938 [00:05<00:00, 173.45it/s, accuracy=0.2620, loss=2.1958]
100%|  | 157/157 [00:00<00:00, 496.47it/s, accuracy=0.3618, loss=2.1098]
Epoch: 17/100



100%|  | 938/938 [00:05<00:00, 174.68it/s, accuracy=0.2980, loss=2.1740]
100%|  | 157/157 [00:00<00:00, 485.76it/s, accuracy=0.4212, loss=2.0757]
Epoch: 18/100



100%|  | 938/938 [00:05<00:00, 180.39it/s, accuracy=0.3351, loss=2.1497]
100%|  | 157/157 [00:00<00:00, 430.39it/s, accuracy=0.4899, loss=2.0368]
Epoch: 19/100



100%|  | 938/938 [00:04<00:00, 198.03it/s, accuracy=0.3612, loss=2.1275]
100%|  | 157/157 [00:00<00:00, 451.10it/s, accuracy=0.5266, loss=1.9997]
Epoch: 20/100



100%|  | 938/938 [00:04<00:00, 199.45it/s, accuracy=0.3999, loss=2.1028]
100%|  | 157/157 [00:00<00:00, 428.73it/s, accuracy=0.5783, loss=1.9622]
Epoch: 21/100



100%|  | 938/938 [00:04<00:00, 194.05it/s, accuracy=0.4158, loss=2.0836]
100%|  | 157/157 [00:00<00:00, 443.94it/s, accuracy=0.5879, loss=1.9383]
Epoch: 22/100



100%|  | 938/938 [00:05<00:00, 185.39it/s, accuracy=0.4271, loss=2.0657]
100%|  | 157/157 [00:00<00:00, 530.34it/s, accuracy=0.5999, loss=1.9165]
Epoch: 23/100



100%|  | 938/938 [00:05<00:00, 180.34it/s, accuracy=0.4388, loss=2.0502]
100%|  | 157/157 [00:00<00:00, 475.98it/s, accuracy=0.6068, loss=1.8986]
Epoch: 24/100


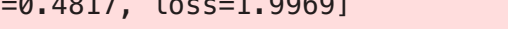
100%|  | 938/938 [00:04<00:00, 196.44it/s, accuracy=0.4434, loss=2.0401]
100%|  | 157/157 [00:00<00:00, 486.52it/s, accuracy=0.6091, loss=1.8889]
Epoch: 25/100


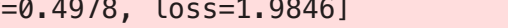
100%|  | 938/938 [00:04<00:00, 194.00it/s, accuracy=0.4518, loss=2.0298]
100%|  | 157/157 [00:00<00:00, 461.56it/s, accuracy=0.6118, loss=1.8812]
Epoch: 26/100


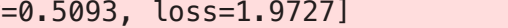
100%|  | 938/938 [00:04<00:00, 195.07it/s, accuracy=0.4587, loss=2.0206]
100%|  | 157/157 [00:00<00:00, 477.50it/s, accuracy=0.6147, loss=1.8743]
Epoch: 27/100



100%|  | 938/938 [00:04<00:00, 196.63it/s, accuracy=0.4667, loss=2.0118]
100%|  | 157/157 [00:00<00:00, 460.50it/s, accuracy=0.6169, loss=1.8663]
Epoch: 28/100



100%|  | 938/938 [00:04<00:00, 194.21it/s, accuracy=0.4717, loss=2.0050]
100%|  | 157/157 [00:00<00:00, 473.61it/s, accuracy=0.6262, loss=1.8528]
Epoch: 29/100



100%|  | 938/938 [00:04<00:00, 195.23it/s, accuracy=0.4817, loss=1.9969]
100%|  | 157/157 [00:00<00:00, 462.92it/s, accuracy=0.6718, loss=1.8284]
Epoch: 30/100



100%|  | 938/938 [00:05<00:00, 183.27it/s, accuracy=0.4978, loss=1.9846]
100%|  | 157/157 [00:00<00:00, 516.37it/s, accuracy=0.6953, loss=1.8108]
Epoch: 31/100



100%|  | 938/938 [00:05<00:00, 185.92it/s, accuracy=0.5093, loss=1.9727]
100%|  | 157/157 [00:00<00:00, 391.81it/s, accuracy=0.7032, loss=1.7988]
Epoch: 32/100



100%|  | 938/938 [00:05<00:00, 183.22it/s, accuracy=0.5202, loss=1.9617]
100%|  | 157/157 [00:00<00:00, 464.33it/s, accuracy=0.7100, loss=1.7882]
Epoch: 33/100



100%|  | 938/938 [00:04<00:00, 194.47it/s, accuracy=0.5333, loss=1.9494]
100%|  | 157/157 [00:00<00:00, 410.10it/s, accuracy=0.7089, loss=1.7856]
Epoch: 34/100


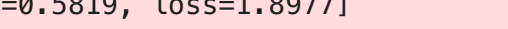
100%|  | 938/938 [00:05<00:00, 187.35it/s, accuracy=0.5436, loss=1.9383]
100%|  | 157/157 [00:00<00:00, 475.99it/s, accuracy=0.7160, loss=1.7776]
Epoch: 35/100


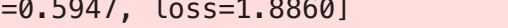
100%|  | 938/938 [00:05<00:00, 185.57it/s, accuracy=0.5521, loss=1.9302]
100%|  | 157/157 [00:00<00:00, 492.62it/s, accuracy=0.7181, loss=1.7746]
Epoch: 36/100


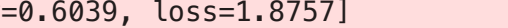
100%|  | 938/938 [00:05<00:00, 181.71it/s, accuracy=0.5609, loss=1.9198]
100%|  | 157/157 [00:00<00:00, 426.56it/s, accuracy=0.7189, loss=1.7705]
Epoch: 37/100



100%|  | 938/938 [00:04<00:00, 191.37it/s, accuracy=0.5689, loss=1.9126]
100%|  | 157/157 [00:00<00:00, 442.01it/s, accuracy=0.7217, loss=1.7678]
Epoch: 38/100



100%|  | 938/938 [00:04<00:00, 195.17it/s, accuracy=0.5777, loss=1.9032]
100%|  | 157/157 [00:00<00:00, 439.42it/s, accuracy=0.7264, loss=1.7619]
Epoch: 39/100



100%|  | 938/938 [00:04<00:00, 196.53it/s, accuracy=0.5819, loss=1.8977]
100%|  | 157/157 [00:00<00:00, 442.62it/s, accuracy=0.7300, loss=1.7579]
Epoch: 40/100



100%|  | 938/938 [00:04<00:00, 194.06it/s, accuracy=0.5947, loss=1.8860]
100%|  | 157/157 [00:00<00:00, 444.47it/s, accuracy=0.7351, loss=1.7533]
Epoch: 41/100



100%|  | 938/938 [00:04<00:00, 194.97it/s, accuracy=0.6039, loss=1.8757]
100%|  | 157/157 [00:00<00:00, 472.68it/s, accuracy=0.7398, loss=1.7473]
Epoch: 42/100



100%|  | 938/938 [00:04<00:00, 191.76it/s, accuracy=0.6109, loss=1.8685]
100%|  | 157/157 [00:00<00:00, 463.34it/s, accuracy=0.7371, loss=1.7479]
Epoch: 43/100



100%|  | 938/938 [00:04<00:00, 191.56it/s, accuracy=0.6221, loss=1.8589]
100%|  | 157/157 [00:00<00:00, 471.25it/s, accuracy=0.7438, loss=1.7419]
Epoch: 44/100


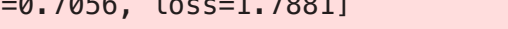
100%|  | 938/938 [00:04<00:00, 192.71it/s, accuracy=0.6294, loss=1.8499]
100%|  | 157/157 [00:00<00:00, 456.62it/s, accuracy=0.7471, loss=1.7305]
Epoch: 45/100


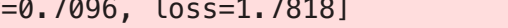
100%|  | 938/938 [00:04<00:00, 193.38it/s, accuracy=0.6612, loss=1.8335]
100%|  | 157/157 [00:00<00:00, 461.05it/s, accuracy=0.8135, loss=1.7036]
Epoch: 46/100


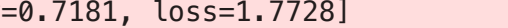
100%|  | 938/938 [00:04<00:00, 194.82it/s, accuracy=0.6784, loss=1.8189]
100%|  | 157/157 [00:00<00:00, 470.18it/s, accuracy=0.8199, loss=1.6927]
Epoch: 47/100



100%|  | 938/938 [00:04<00:00, 192.85it/s, accuracy=0.6852, loss=1.8095]
100%|  | 157/157 [00:00<00:00, 469.21it/s, accuracy=0.8174, loss=1.6875]
Epoch: 48/100



100%|  | 938/938 [00:04<00:00, 192.17it/s, accuracy=0.6968, loss=1.7971]
100%|  | 157/157 [00:00<00:00, 459.38it/s, accuracy=0.8216, loss=1.6799]
Epoch: 49/100



100%|  | 938/938 [00:05<00:00, 183.41it/s, accuracy=0.7056, loss=1.7881]
100%|  | 157/157 [00:00<00:00, 435.21it/s, accuracy=0.8267, loss=1.6721]
Epoch: 50/100



100%|  | 938/938 [00:04<00:00, 194.99it/s, accuracy=0.7096, loss=1.7818]
100%|  | 157/157 [00:00<00:00, 436.39it/s, accuracy=0.8298, loss=1.6679]
Epoch: 51/100



100%|  | 938/938 [00:04<00:00, 195.21it/s, accuracy=0.7181, loss=1.7728]
100%|  | 157/157 [00:00<00:00, 428.96it/s, accuracy=0.8304, loss=1.6629]
Epoch: 52/100



100%|  | 938/938 [00:04<00:00, 194.91it/s, accuracy=0.7228, loss=1.7652]
100%|  | 157/157 [00:00<00:00, 439.41it/s, accuracy=0.8328, loss=1.6579]
Epoch: 53/100



100%|  | 938/938 [00:04<00:00, 194.73it/s, accuracy=0.7285, loss=1.7607]
100%|  | 157/157 [00:00<00:00, 427.10it/s, accuracy=0.8369, loss=1.6539]
Epoch: 54/100


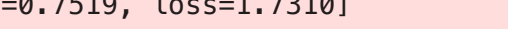
100%|  | 938/938 [00:04<00:00, 191.41it/s, accuracy=0.7333, loss=1.7529]
100%|  | 157/157 [00:00<00:00, 470.92it/s, accuracy=0.8331, loss=1.6543]
Epoch: 55/100


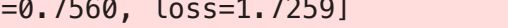
100%|  | 938/938 [00:04<00:00, 192.86it/s, accuracy=0.7371, loss=1.7488]
100%|  | 157/157 [00:00<00:00, 472.80it/s, accuracy=0.8373, loss=1.6493]
Epoch: 56/100


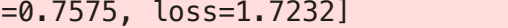
100%|  | 938/938 [00:04<00:00, 193.67it/s, accuracy=0.7403, loss=1.7446]
100%|  | 157/157 [00:00<00:00, 476.44it/s, accuracy=0.8382, loss=1.6452]
Epoch: 57/100



100%|  | 938/938 [00:04<00:00, 194.03it/s, accuracy=0.7464, loss=1.7382]
100%|  | 157/157 [00:00<00:00, 464.09it/s, accuracy=0.8420, loss=1.6401]
Epoch: 58/100



100%|  | 938/938 [00:04<00:00, 192.87it/s, accuracy=0.7493, loss=1.7341]
100%|  | 157/157 [00:00<00:00, 460.41it/s, accuracy=0.8429, loss=1.6396]
Epoch: 59/100



100%|  | 938/938 [00:04<00:00, 192.57it/s, accuracy=0.7519, loss=1.7310]
100%|  | 157/157 [00:00<00:00, 469.92it/s, accuracy=0.8433, loss=1.6384]
Epoch: 60/100



100%|  | 938/938 [00:04<00:00, 194.03it/s, accuracy=0.7560, loss=1.7259]
100%|  | 157/157 [00:00<00:00, 469.44it/s, accuracy=0.8466, loss=1.6348]
Epoch: 61/100



100%|  | 938/938 [00:04<00:00, 192.99it/s, accuracy=0.7575, loss=1.7232]
100%|  | 157/157 [00:00<00:00, 471.72it/s, accuracy=0.8460, loss=1.6335]
Epoch: 62/100



100%|  | 938/938 [00:04<00:00, 193.37it/s, accuracy=0.7619, loss=1.7184]
100%|  | 157/157 [00:00<00:00, 463.57it/s, accuracy=0.8462, loss=1.6314]
Epoch: 63/100



100%|  | 938/938 [00:04<00:00, 192.89it/s, accuracy=0.7671, loss=1.7147]
100%|  | 157/157 [00:00<00:00, 466.09it/s, accuracy=0.8495, loss=1.6298]
Epoch: 64/100


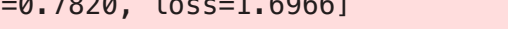
100%|  | 938/938 [00:04<00:00, 192.37it/s, accuracy=0.7688, loss=1.7115]
100%|  | 157/157 [00:00<00:00, 462.45it/s, accuracy=0.8511, loss=1.6280]
Epoch: 65/100


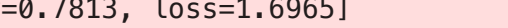
100%|  | 938/938 [00:04<00:00, 194.37it/s, accuracy=0.7689, loss=1.7091]
100%|  | 157/157 [00:00<00:00, 469.05it/s, accuracy=0.8513, loss=1.6257]
Epoch: 66/100


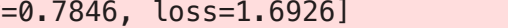
100%|  | 938/938 [00:04<00:00, 192.46it/s, accuracy=0.7735, loss=1.7055]
100%|  | 157/157 [00:00<00:00, 474.32it/s, accuracy=0.8479, loss=1.6271]
Epoch: 67/100



100%|  | 938/938 [00:04<00:00, 193.48it/s, accuracy=0.7754, loss=1.7029]
100%|  | 157/157 [00:00<00:00, 467.79it/s, accuracy=0.8539, loss=1.6220]
Epoch: 68/100



100%|  | 938/938 [00:04<00:00, 194.06it/s, accuracy=0.7783, loss=1.7000]
100%|  | 157/157 [00:00<00:00, 474.30it/s, accuracy=0.8550, loss=1.6210]
Epoch: 69/100



100%|  | 938/938 [00:04<00:00, 194.16it/s, accuracy=0.7820, loss=1.6966]
100%|  | 157/157 [00:00<00:00, 465.09it/s, accuracy=0.8547, loss=1.6208]
Epoch: 70/100



100%|  | 938/938 [00:04<00:00, 192.72it/s, accuracy=0.7813, loss=1.6965]
100%|  | 157/157 [00:00<00:00, 469.79it/s, accuracy=0.8556, loss=1.6193]
Epoch: 71/100



100%|  | 938/938 [00:04<00:00, 192.85it/s, accuracy=0.7846, loss=1.6926]
100%|  | 157/157 [00:00<00:00, 449.96it/s, accuracy=0.8575, loss=1.6168]
Epoch: 72/100



100%|  | 938/938 [00:04<00:00, 193.14it/s, accuracy=0.7873, loss=1.6896]
100%|  | 157/157 [00:00<00:00, 467.66it/s, accuracy=0.8570, loss=1.6162]
Epoch: 73/100



100%|  | 938/938 [00:05<00:00, 183.05it/s, accuracy=0.7883, loss=1.6881]
100%|  | 157/157 [00:00<00:00, 439.08it/s, accuracy=0.8579, loss=1.6146]
Epoch: 74/100


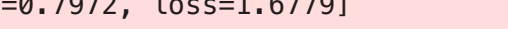
100%|  | 938/938 [00:04<00:00, 193.56it/s, accuracy=0.7882, loss=1.6877]
100%|  | 157/157 [00:00<00:00, 429.35it/s, accuracy=0.8573, loss=1.6150]
Epoch: 75/100


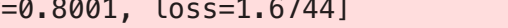
100%|  | 938/938 [00:04<00:00, 193.46it/s, accuracy=0.7906, loss=1.6854]
100%|  | 157/157 [00:00<00:00, 432.93it/s, accuracy=0.8596, loss=1.6129]
Epoch: 76/100


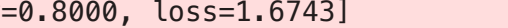
100%|  | 938/938 [00:04<00:00, 193.17it/s, accuracy=0.7917, loss=1.6833]
100%|  | 157/157 [00:00<00:00, 457.26it/s, accuracy=0.8591, loss=1.6129]
Epoch: 77/100



100%|  | 938/938 [00:04<00:00, 192.22it/s, accuracy=0.7944, loss=1.6807]
100%|  | 157/157 [00:00<00:00, 468.97it/s, accuracy=0.8628, loss=1.6109]
Epoch: 78/100



100%|  | 938/938 [00:04<00:00, 192.39it/s, accuracy=0.7990, loss=1.6764]
100%|  | 157/157 [00:00<00:00, 420.09it/s, accuracy=0.8612, loss=1.6099]
Epoch: 79/100



100%|  | 938/938 [00:04<00:00, 192.49it/s, accuracy=0.7972, loss=1.6779]
100%|  | 157/157 [00:00<00:00, 467.70it/s, accuracy=0.8607, loss=1.6113]
Epoch: 80/100



100%|  | 938/938 [00:04<00:00, 190.30it/s, accuracy=0.8001, loss=1.6744]
100%|  | 157/157 [00:00<00:00, 472.07it/s, accuracy=0.8644, loss=1.6083]
Epoch: 81/100



100%|  | 938/938 [00:04<00:00, 193.13it/s, accuracy=0.8000, loss=1.6743]
100%|  | 157/157 [00:00<00:00, 456.24it/s, accuracy=0.8623, loss=1.6084]
Epoch: 82/100



100%|  | 938/938 [00:04<00:00, 192.58it/s, accuracy=0.8018, loss=1.6726]
100%|  | 157/157 [00:00<00:00, 470.32it/s, accuracy=0.8609, loss=1.6098]
Epoch: 83/100



100%|  | 938/938 [00:04<00:00, 193.15it/s, accuracy=0.8038, loss=1.6702]
100%|  | 157/157 [00:00<00:00, 468.85it/s, accuracy=0.8620, loss=1.6078]
Epoch: 84/100


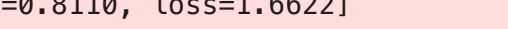
100%|  | 938/938 [00:04<00:00, 192.42it/s, accuracy=0.8043, loss=1.6697]
100%|  | 157/157 [00:00<00:00, 472.73it/s, accuracy=0.8627, loss=1.6068]
Epoch: 85/100


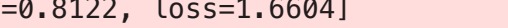
100%|  | 938/938 [00:04<00:00, 192.40it/s, accuracy=0.8065, loss=1.6667]
100%|  | 157/157 [00:00<00:00, 455.30it/s, accuracy=0.8629, loss=1.6063]
Epoch: 86/100


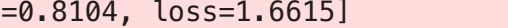
100%|  | 938/938 [00:04<00:00, 192.12it/s, accuracy=0.8069, loss=1.6664]
100%|  | 157/157 [00:00<00:00, 469.48it/s, accuracy=0.8643, loss=1.6045]
Epoch: 87/100

100%|  | 938/938 [00:04<00:00, 192.94it/s, accuracy=0.8093, loss=1.6638]
100%|  | 157/157 [00:00<00:00, 461.51it/s, accuracy=0.8647, loss=1.6048]
Epoch: 88/100

100%|  | 938/938 [00:04<00:00, 192.63it/s, accuracy=0.8102, loss=1.6629]
100%|  | 157/157 [00:00<00:00, 469.32it/s, accuracy=0.8653, loss=1.6030]
Epoch: 89/100

100%|  | 938/938 [00:04<00:00, 192.89it/s, accuracy=0.8110, loss=1.6622]
100%|  | 157/157 [00:00<00:00, 466.72it/s, accuracy=0.8653, loss=1.6028]
Epoch: 90/100

100%|  | 938/938 [00:05<00:00, 179.55it/s, accuracy=0.8122, loss=1.6604]
100%|  | 157/157 [00:00<00:00, 471.74it/s, accuracy=0.8652, loss=1.6029]
Epoch: 91/100

100%|  | 938/938 [00:05<00:00, 185.70it/s, accuracy=0.8104, loss=1.6615]
100%|  | 157/157 [00:00<00:00, 471.59it/s, accuracy=0.8648, loss=1.6020]
Epoch: 92/100

```
headers = history2.keys()
rows = zip(*history2.values())
print(tabulate(rows, headers=headers, tablefmt="pretty"))
```


loss_train	acc_train	loss_test	acc_test
2.3023917994181313	0.11323333333333334	2.3015937885284425	0.1276
2.301779421488444	0.12786666666666666	2.3004300525665284	0.1411
2.301105199686686	0.12621666666666667	2.299012487792969	0.1157
2.300234745279948	0.10966666666666666	2.297085223007202	0.0997
2.2989735811869303	0.0999	2.2939566635131836	0.098
2.2967103847503663	0.09873333333333334	2.2872416290283204	0.098
2.291719249725342	0.09871666666666666	2.2726911727905272	0.098
2.2843660423278807	0.10476666666666666	2.2576027240753174	0.1703
2.277094981002808	0.14411666666666667	2.2443132095336913	0.2138
2.269759626261393	0.17028333333333334	2.233527021026611	0.2413
2.2631625747680664	0.18615	2.224640707397461	0.2526
2.2553498407999673	0.19888333333333333	2.2138193153381347	0.2598
2.246927773920695	0.21073333333333333	2.197453458404541	0.262
2.2338521931966144	0.22133333333333333	2.1702111965179443	0.2765
2.2170833276112876	0.23645	2.141159670257568	0.3091
2.195816873041789	0.26201666666666667	2.10979799156189	0.3618
2.1739661933898926	0.29801666666666665	2.0757017070770263	0.4212
2.149686218770345	0.33513333333333334	2.036761261367798	0.4899
2.1274956385294597	0.36121666666666667	1.9996891317367553	0.5266
2.1027778367360432	0.39986666666666665	1.9622000730514526	0.5783
2.08357397702535	0.41575	1.9382757385253906	0.5879
2.0657051109313964	0.4271	1.9164779043197633	0.5999
2.0502357849121093	0.4388	1.8985953496932984	0.6068
2.0400845144907636	0.44336666666666667	1.8888752071380615	0.6091
2.029832869974772	0.45183333333333333	1.881192015838623	0.6118

2.020633637491862	0.4586833333333333	1.8742718015670776	0.6147
2.0117791975657147	0.46675	1.8662612926483155	0.6169
2.004992022387187	0.4717166666666667	1.852794076538086	0.6262
1.9968504940032958	0.4816666666666667	1.8283619041442871	0.6718
1.9846370325724283	0.4977666666666667	1.810757007408142	0.6953
1.9726618200937907	0.5093166666666666	1.7988089021682738	0.7032
1.9616771308898926	0.5202333333333333	1.788227031135559	0.71
1.9494058584213256	0.5333166666666667	1.785608842086792	0.7089
1.938273152669271	0.54365	1.7775830883026122	0.716
1.9302074719111124	0.5521166666666667	1.7745666666030884	0.7181
1.9197937030792236	0.5609	1.770523743247986	0.7189
1.912576283899943	0.56895	1.767826984024048	0.7217
1.9031891501108806	0.5777333333333333	1.7619387874603272	0.7264
1.8976596705118816	0.5819333333333333	1.7578776252746582	0.73
1.8859897172927858	0.5946666666666667	1.7533020643234254	0.7351
1.8756669408798219	0.6038833333333333	1.7473428983688355	0.7398
1.8685280645370483	0.6108666666666667	1.7478581304550171	0.7371
1.8588582169850667	0.6220833333333333	1.741906698036194	0.7438
1.8498946297963461	0.62945	1.7304668962478638	0.7471
1.8334701229731243	0.6612166666666667	1.7035838775634766	0.8135
1.8189258522669474	0.6784	1.6926774543762206	0.8199
1.8095341126124065	0.6851666666666667	1.6874780223846435	0.8174
1.797077129236857	0.6968166666666666	1.6799101720809937	0.8216
1.78806405321757	0.7055833333333333	1.6721400651931764	0.8267
1.7817921817143758	0.7095666666666667	1.6679342073440553	0.8298
1.7728199191411336	0.7180666666666666	1.6629396572113038	0.8304

1.7651949092229207	0.7228166666666667	1.6579278593063353	0.8328
1.760677334022522	0.7285166666666667	1.6539058145523071	0.8369
1.7529308547973632	0.7333	1.6542914155960082	0.8331
1.7487792910893758	0.7370833333333333	1.6492848066329957	0.8373
1.7446007921854656	0.7403166666666666	1.6452130805969238	0.8382
1.7381847756067912	0.7464333333333333	1.6401126943588258	0.842
1.7340990405400594	0.7493166666666666	1.6395992765426635	0.8429
1.731024103609721	0.7519166666666667	1.6384345928192139	0.8433
1.7259252298355103	0.75605	1.6348146808624267	0.8466
1.7231981597264607	0.7575	1.6335084999084473	0.846
1.7184376949946085	0.76185	1.6313756980895997	0.8462
1.714709228960673	0.7671	1.629794972229004	0.8495
1.7114928155263265	0.7688333333333334	1.628007110595703	0.8511
1.7090858845392862	0.7688666666666667	1.6257491724014281	0.8513
1.7055176521937052	0.7734833333333333	1.6271148765563965	0.8479
1.7028684744517009	0.7754166666666666	1.6220443759918213	0.8539
1.7000056770324707	0.7783	1.620961150741577	0.855
1.6965951473236085	0.7820333333333334	1.6208162227630616	0.8547
1.6964880917867025	0.78135	1.6193018367767333	0.8556
1.6926296309789022	0.7845666666666666	1.6168323093414307	0.8575
1.6896281602223715	0.78725	1.6162059545516967	0.857
1.6881227117538453	0.78825	1.6146186946868897	0.8579
1.6876972617467245	0.7881666666666667	1.6150481632232665	0.8573
1.6853521058400471	0.7906166666666666	1.612942784500122	0.8596
1.6833227584203085	0.7917	1.6128609712600708	0.8591
1.6806547307332356	0.79435	1.610928844833374	0.8628
1.6763857836405436	0.7990333333333334	1.6098757152557372	0.8612

1.6779367630004882	0.7972333333333333	1.611296816444397	0.8607
1.6743898506800334	0.8001333333333334	1.6082801803588866	0.8644
1.6742977463404338	0.8	1.608447886276245	0.8623
1.672567369969686	0.8017666666666666	1.6097805156707763	0.8609
1.6701994068145751	0.8038	1.607767364692688	0.862
1.6696795998891194	0.8042833333333334	1.6067954328536986	0.8627
1.6666988358179728	0.8064833333333333	1.6063465614318848	0.8629
1.6664206731796265	0.8069166666666666	1.6045262624740602	0.8643
1.6637693227767945	0.80925	1.60480203666687	0.8647
1.662858319536845	0.8101666666666667	1.6029749931335449	0.8653
1.6622229460398357	0.811	1.602805283164978	0.8653
1.660368582979838	0.8122166666666667	1.6029315828323365	0.8652
1.661468242963155	0.8104166666666667	1.6020055009841918	0.8648
1.6579865900675457	0.8142666666666667	1.6009479763031005	0.8667
1.6568754144032796	0.8149333333333333	1.5998276706695556	0.8674
1.655471849822998	0.8173166666666667	1.6014277215957642	0.8663
1.6553486333847045	0.81685	1.6007486192703246	0.8659
1.6538977573394775	0.8179833333333333	1.599662559890747	0.8681
1.6531105925242107	0.8177666666666666	1.598543883895874	0.8677
1.6526153662363687	0.81885	1.5986144170761107	0.8668
1.652674685160319	0.8181666666666667	1.5974585599899291	0.8691
1.6499494305928548	0.8217333333333333	1.602588733291626	0.866

```
In [22]: epochs = range(1, len(history2["loss_train"]) + 1)
          draw_loss_test(epochs, history2)
```



```
In [26]: net, history = training(
    net,
    train_loader,
    test_loader,
    loss_fn,
    metric,
    optimizer,
    update_period=5,
    epoch_max=100,
    device=device,
    early_stopping_accuracy=0.98
)
```

```
100%|██████████████████████████████████████████████████████████████████████████| 938/938  
[00:03<00:00, 246.67it/s]  
Train loss: 2.3023  
Train accuracy: 0.1151
```


```
100%|██████████████████████████████████████████████████████████████████████████| 157/157  
[00:00<00:00, 713.15it/s]  
Test loss: 2.3015  
Test accuracy: 0.1905
```

```
100%|██████████████████████████████████████████████████████████████████████████| 938/938  
[00:03<00:00, 250.51it/s]  
Train loss: 2.3016  
Train accuracy: 0.1659
```

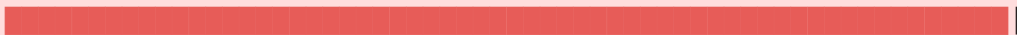
```
100%|██████████████████████████████████████████████████████████████████████████| 157/157  
[00:00<00:00, 732.43it/s]  
Test loss: 2.3003  
Test accuracy: 0.2611
```


[illegible]

Train loss: 2.3008
Train accuracy: 0.1692

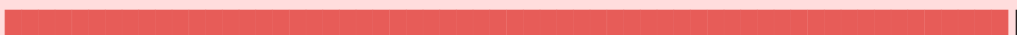
100%|  | 157/157
[00:00<00:00, 637.85it/s]
Test loss: 2.2987
Test accuracy: 0.2170


Epoka: 4/100

100%|  | 938/938
[00:03<00:00, 249.93it/s]
Train loss: 2.2998
Train accuracy: 0.1371

100%|  | 157/157
[00:00<00:00, 736.85it/s]
Test loss: 2.2966
Test accuracy: 0.1370


Epoka: 5/100

100%|  | 938/938
[00:03<00:00, 249.55it/s]
Train loss: 2.2984
Train accuracy: 0.1083

100%|  | 157/157
[00:00<00:00, 719.84it/s]
Test loss: 2.2933
Test accuracy: 0.1096

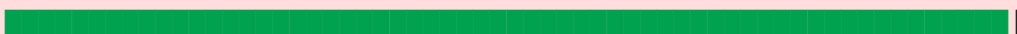
Epoka: 6/100

100%|  | 938/938
[00:03<00:00, 247.41it/s]
Train loss: 2.2958
Train accuracy: 0.1007

100%|  | 157/157
[00:00<00:00, 715.40it/s]
Test loss: 2.2862
Test accuracy: 0.0988

Epoka: 7/100


100%|  | 938/938
[00:03<00:00, 247.97it/s]
Train loss: 2.2903
Train accuracy: 0.0990

100%|  | 157/157
[00:00<00:00, 724.32it/s]
Test loss: 2.2714
Test accuracy: 0.0989


Epoka: 8/100

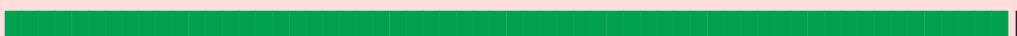
100%|  | 938/938
[00:03<00:00, 247.94it/s]

Train loss: 2.2824
Train accuracy: 0.1140

100%|  | 157/157
[00:00<00:00, 721.41it/s]
Test loss: 2.2551
Test accuracy: 0.1885

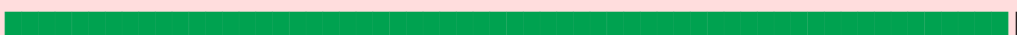
Epoka: 9/100

100%|  | 938/938
[00:03<00:00, 246.45it/s]
Train loss: 2.2745
Train accuracy: 0.1557

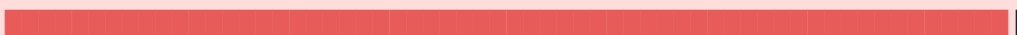
100%|  | 157/157
[00:00<00:00, 713.42it/s]
Test loss: 2.2403
Test accuracy: 0.2304

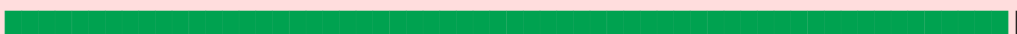
Epoka: 10/100

100%|  | 938/938
[00:03<00:00, 248.26it/s]
Train loss: 2.2661
Train accuracy: 0.1824

100%|  | 157/157
[00:00<00:00, 690.45it/s]
Test loss: 2.2268
Test accuracy: 0.2549

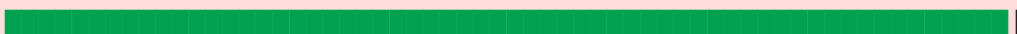
Epoka: 11/100

100%|  | 938/938
[00:03<00:00, 247.61it/s]
Train loss: 2.2579
Train accuracy: 0.2041

100%|  | 157/157
[00:00<00:00, 685.75it/s]
Test loss: 2.2118
Test accuracy: 0.2700

Epoka: 12/100


100%|  | 938/938
[00:03<00:00, 244.63it/s]
Train loss: 2.2462
Train accuracy: 0.2210

100%|  | 157/157
[00:00<00:00, 685.65it/s]
Test loss: 2.1881
Test accuracy: 0.2797


Epoka: 13/100


100%|  | 938/938
[00:03<00:00, 246.40it/s]

Train loss: 2.2321
Train accuracy: 0.2276


100%|  | 157/157
[00:00<00:00, 686.89it/s]
Test loss: 2.1604
Test accuracy: 0.2811

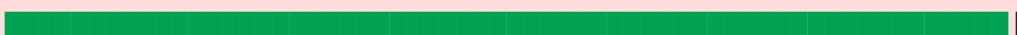
Epoka: 14/100

100%|  | 938/938
[00:03<00:00, 246.54it/s]
Train loss: 2.2143
Train accuracy: 0.2364

100%|  | 157/157
[00:00<00:00, 684.99it/s]
Test loss: 2.1322
Test accuracy: 0.3105


Epoka: 15/100

100%|  | 938/938
[00:03<00:00, 245.21it/s]
Train loss: 2.1950
Train accuracy: 0.2673


100%|  | 157/157
[00:00<00:00, 687.70it/s]
Test loss: 2.1030
Test accuracy: 0.3761

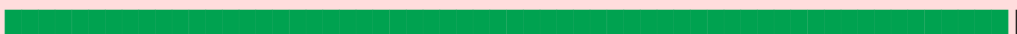
Epoka: 16/100

100%|  | 938/938
[00:04<00:00, 230.16it/s]
Train loss: 2.1716
Train accuracy: 0.3055

100%|  | 157/157
[00:00<00:00, 699.42it/s]
Test loss: 2.0693
Test accuracy: 0.4537

Epoka: 17/100


100%|  | 938/938
[00:03<00:00, 244.60it/s]
Train loss: 2.1495
Train accuracy: 0.3418

100%|  | 157/157
[00:00<00:00, 702.59it/s]
Test loss: 2.0328
Test accuracy: 0.5003


Epoka: 18/100

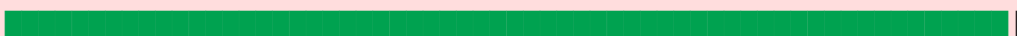
100%|  | 938/938
[00:03<00:00, 241.76it/s]

Train loss: 2.1257
Train accuracy: 0.3679


100%|  | 157/157
[00:00<00:00, 704.33it/s]
Test loss: 1.9981
Test accuracy: 0.5231

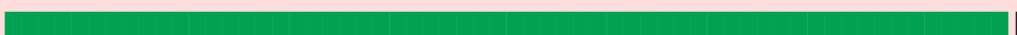
Epoka: 19/100

100%|  | 938/938
[00:03<00:00, 245.17it/s]
Train loss: 2.1079
Train accuracy: 0.3816

100%|  | 157/157
[00:00<00:00, 698.16it/s]
Test loss: 1.9725
Test accuracy: 0.5324


Epoka: 20/100

100%|  | 938/938
[00:03<00:00, 244.77it/s]
Train loss: 2.0895
Train accuracy: 0.3948

100%|  | 157/157
[00:00<00:00, 630.84it/s]
Test loss: 1.9506
Test accuracy: 0.5425

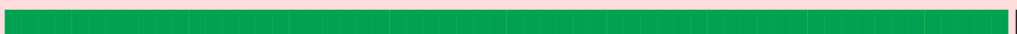
Epoka: 21/100

100%|  | 938/938
[00:03<00:00, 244.56it/s]
Train loss: 2.0776
Train accuracy: 0.4012

100%|  | 157/157
[00:00<00:00, 698.37it/s]
Test loss: 1.9338
Test accuracy: 0.5478

Epoka: 22/100


100%|  | 938/938
[00:03<00:00, 242.25it/s]
Train loss: 2.0644
Train accuracy: 0.4132

100%|  | 157/157
[00:00<00:00, 679.73it/s]
Test loss: 1.9105
Test accuracy: 0.5792


Epoka: 23/100

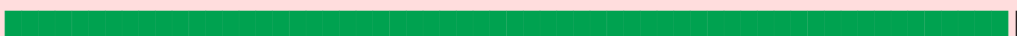
100%|  | 938/938
[00:04<00:00, 208.82it/s]

Train loss: 2.0518
Train accuracy: 0.4268

100%|  | 157/157
[00:00<00:00, 689.74it/s]
Test loss: 1.8863
Test accuracy: 0.6171

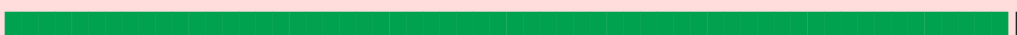
Epoka: 24/100

100%|  | 938/938
[00:04<00:00, 233.79it/s]
Train loss: 2.0395
Train accuracy: 0.4378

100%|  | 157/157
[00:00<00:00, 614.96it/s]
Test loss: 1.8687
Test accuracy: 0.6237


Epoka: 25/100

100%|  | 938/938
[00:04<00:00, 231.02it/s]
Train loss: 2.0272
Train accuracy: 0.4524

100%|  | 157/157
[00:00<00:00, 691.65it/s]
Test loss: 1.8454
Test accuracy: 0.6784

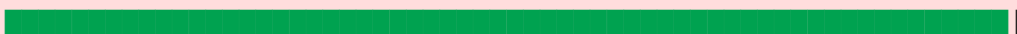
Epoka: 26/100

100%|  | 938/938
[00:04<00:00, 230.70it/s]
Train loss: 2.0128
Train accuracy: 0.4710

100%|  | 157/157
[00:00<00:00, 682.00it/s]
Test loss: 1.8311
Test accuracy: 0.6865

Epoka: 27/100


100%|  | 938/938
[00:03<00:00, 238.10it/s]
Train loss: 2.0013
Train accuracy: 0.4828

100%|  | 157/157
[00:00<00:00, 621.07it/s]
Test loss: 1.8200
Test accuracy: 0.6917


Epoka: 28/100

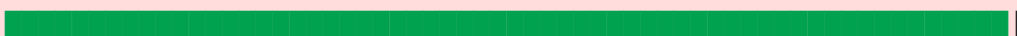
100%|  | 938/938
[00:03<00:00, 242.90it/s]

Train loss: 1.9901
Train accuracy: 0.4942

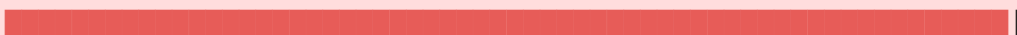
100%|  | 157/157
[00:00<00:00, 705.30it/s]
Test loss: 1.8102
Test accuracy: 0.6980

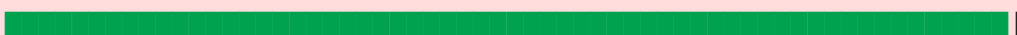
Epoka: 29/100

100%|  | 938/938
[00:03<00:00, 243.15it/s]
Train loss: 1.9806
Train accuracy: 0.5027

100%|  | 157/157
[00:00<00:00, 633.52it/s]
Test loss: 1.8058
Test accuracy: 0.6989

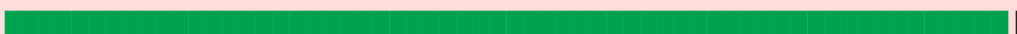
Epoka: 30/100

100%|  | 938/938
[00:04<00:00, 220.69it/s]
Train loss: 1.9697
Train accuracy: 0.5134

100%|  | 157/157
[00:00<00:00, 665.14it/s]
Test loss: 1.7973
Test accuracy: 0.7030

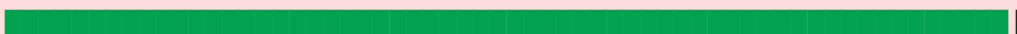
Epoka: 31/100

100%|  | 938/938
[00:03<00:00, 240.05it/s]
Train loss: 1.9601
Train accuracy: 0.5215

100%|  | 157/157
[00:00<00:00, 619.20it/s]
Test loss: 1.7918
Test accuracy: 0.7092

Epoka: 32/100


100%|  | 938/938
[00:03<00:00, 241.88it/s]
Train loss: 1.9525
Train accuracy: 0.5281

100%|  | 157/157
[00:00<00:00, 683.34it/s]
Test loss: 1.7848
Test accuracy: 0.7122


Epoka: 33/100


100%|  | 938/938
[00:03<00:00, 241.79it/s]

Train loss: 1.9426
Train accuracy: 0.5383


100%|  | 157/157
[00:00<00:00, 688.53it/s]
Test loss: 1.7837
Test accuracy: 0.7116

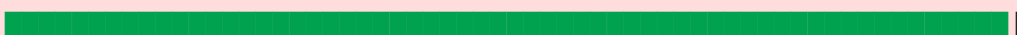
Epoka: 34/100

100%|  | 938/938
[00:03<00:00, 240.03it/s]
Train loss: 1.9323
Train accuracy: 0.5486

100%|  | 157/157
[00:00<00:00, 696.59it/s]
Test loss: 1.7764
Test accuracy: 0.7173

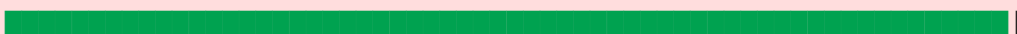
Epoka: 35/100

100%|  | 938/938
[00:03<00:00, 242.13it/s]
Train loss: 1.9245
Train accuracy: 0.5554


100%|  | 157/157
[00:00<00:00, 684.99it/s]
Test loss: 1.7727
Test accuracy: 0.7222

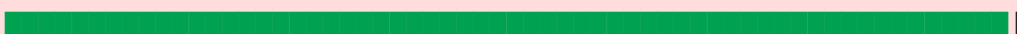
Epoka: 36/100

100%|  | 938/938
[00:03<00:00, 242.05it/s]
Train loss: 1.9151
Train accuracy: 0.5648

100%|  | 157/157
[00:00<00:00, 694.42it/s]
Test loss: 1.7692
Test accuracy: 0.7247

Epoka: 37/100


100%|  | 938/938
[00:03<00:00, 240.56it/s]
Train loss: 1.9070
Train accuracy: 0.5749

100%|  | 157/157
[00:00<00:00, 683.91it/s]
Test loss: 1.7651
Test accuracy: 0.7288


Epoka: 38/100


100%|  | 938/938
[00:03<00:00, 242.10it/s]

Train loss: 1.8974
Train accuracy: 0.5829


100%|  | 157/157
[00:00<00:00, 687.42it/s]
Test loss: 1.7598
Test accuracy: 0.7316

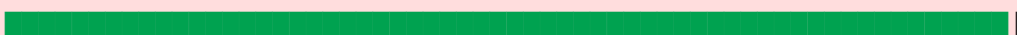
Epoka: 39/100

100%|  | 938/938
[00:03<00:00, 241.95it/s]
Train loss: 1.8892
Train accuracy: 0.5931

100%|  | 157/157
[00:00<00:00, 664.83it/s]
Test loss: 1.7547
Test accuracy: 0.7368

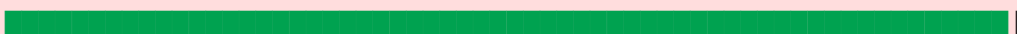
Epoka: 40/100

100%|  | 938/938
[00:03<00:00, 240.57it/s]
Train loss: 1.8795
Train accuracy: 0.6015

100%|  | 157/157
[00:00<00:00, 678.88it/s]
Test loss: 1.7531
Test accuracy: 0.7368

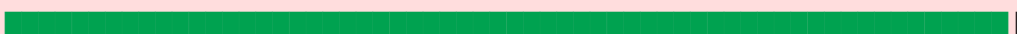
Epoka: 41/100

100%|  | 938/938
[00:03<00:00, 241.56it/s]
Train loss: 1.8692
Train accuracy: 0.6126

100%|  | 157/157
[00:00<00:00, 676.98it/s]
Test loss: 1.7488
Test accuracy: 0.7379

Epoka: 42/100


100%|  | 938/938
[00:03<00:00, 240.14it/s]
Train loss: 1.8629
Train accuracy: 0.6185

100%|  | 157/157
[00:00<00:00, 610.79it/s]
Test loss: 1.7470
Test accuracy: 0.7393


Epoka: 43/100

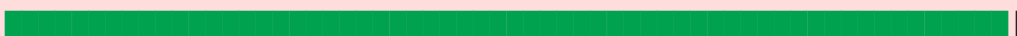
100%|  | 938/938
[00:03<00:00, 241.46it/s]

Train loss: 1.8533
Train accuracy: 0.6288

100%|  | 157/157
[00:00<00:00, 685.25it/s]
Test loss: 1.7453
Test accuracy: 0.7386

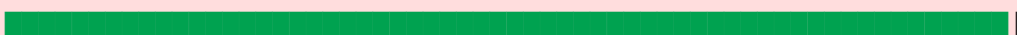
Epoka: 44/100

100%|  | 938/938
[00:03<00:00, 241.30it/s]
Train loss: 1.8466
Train accuracy: 0.6353

100%|  | 157/157
[00:00<00:00, 686.41it/s]
Test loss: 1.7416
Test accuracy: 0.7417


Epoka: 45/100

100%|  | 938/938
[00:03<00:00, 238.09it/s]
Train loss: 1.8416
Train accuracy: 0.6397

100%|  | 157/157
[00:00<00:00, 673.71it/s]
Test loss: 1.7383
Test accuracy: 0.7448

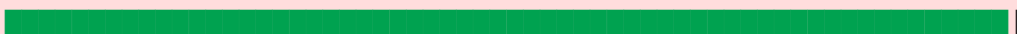
Epoka: 46/100

100%|  | 938/938
[00:04<00:00, 224.69it/s]
Train loss: 1.8341
Train accuracy: 0.6462

100%|  | 157/157
[00:00<00:00, 667.89it/s]
Test loss: 1.7360
Test accuracy: 0.7467

Epoka: 47/100


100%|  | 938/938
[00:03<00:00, 240.58it/s]
Train loss: 1.8300
Train accuracy: 0.6487

100%|  | 157/157
[00:00<00:00, 675.04it/s]
Test loss: 1.7332
Test accuracy: 0.7476

Epoka: 48/100

100%|  | 938/938
[00:03<00:00, 239.23it/s]

Train loss: 1.8217
Train accuracy: 0.6573

100%|  | 157/157

[00:00<00:00, 665.17it/s]

Test loss: 1.7297
Test accuracy: 0.7495

Epoka: 49/100

100%|  | 938/938

[00:03<00:00, 241.38it/s]

Train loss: 1.8162
Train accuracy: 0.6636

100%|  | 157/157

[00:00<00:00, 691.73it/s]

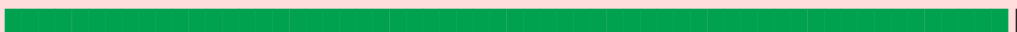
Test loss: 1.7268
Test accuracy: 0.7510

Epoka: 50/100

100%|  | 938/938

[00:03<00:00, 241.39it/s]

Train loss: 1.8118
Train accuracy: 0.6671

100%|  | 157/157

[00:00<00:00, 679.91it/s]

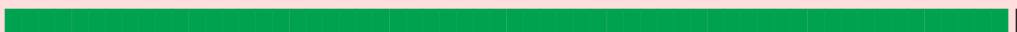
Test loss: 1.7245
Test accuracy: 0.7527

Epoka: 51/100

100%|  | 938/938

[00:03<00:00, 236.94it/s]

Train loss: 1.8076
Train accuracy: 0.6697

100%|  | 157/157

[00:00<00:00, 679.60it/s]

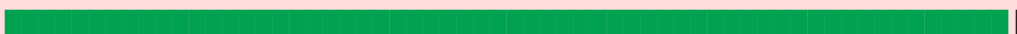
Test loss: 1.7231
Test accuracy: 0.7540

Epoka: 52/100

100%|  | 938/938

[00:03<00:00, 240.62it/s]

Train loss: 1.8018
Train accuracy: 0.6756

100%|  | 157/157

[00:00<00:00, 673.83it/s]


Test loss: 1.7211
Test accuracy: 0.7543

Epoka: 53/100


100%|  | 938/938


[00:03<00:00, 239.81it/s]

Train loss: 1.7996
Train accuracy: 0.6777


100%|  | 157/157
[00:00<00:00, 600.01it/s]
Test loss: 1.7186
Test accuracy: 0.7573

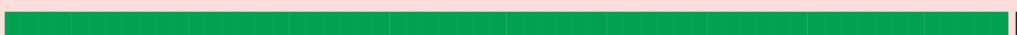
Epoka: 54/100

100%|  | 938/938
[00:03<00:00, 239.60it/s]
Train loss: 1.7941
Train accuracy: 0.6825

100%|  | 157/157
[00:00<00:00, 677.72it/s]
Test loss: 1.7187
Test accuracy: 0.7547


Epoka: 55/100

100%|  | 938/938
[00:03<00:00, 239.97it/s]
Train loss: 1.7922
Train accuracy: 0.6846

100%|  | 157/157
[00:00<00:00, 661.82it/s]
Test loss: 1.7142
Test accuracy: 0.7600

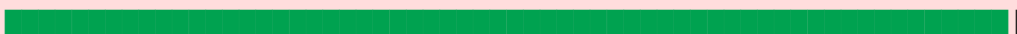
Epoka: 56/100

100%|  | 938/938
[00:03<00:00, 237.85it/s]
Train loss: 1.7889
Train accuracy: 0.6861

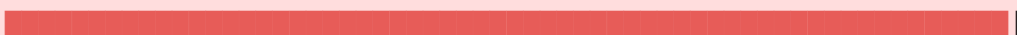
100%|  | 157/157
[00:00<00:00, 665.58it/s]
Test loss: 1.7143
Test accuracy: 0.7592

Epoka: 57/100


100%|  | 938/938
[00:03<00:00, 239.50it/s]
Train loss: 1.7846
Train accuracy: 0.6907

100%|  | 157/157
[00:00<00:00, 658.31it/s]
Test loss: 1.7111
Test accuracy: 0.7620

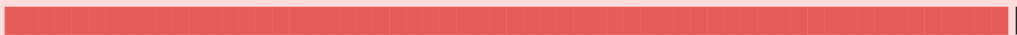
Epoka: 58/100


100%|  | 938/938
[00:03<00:00, 240.21it/s]

Train loss: 1.7822
Train accuracy: 0.6927

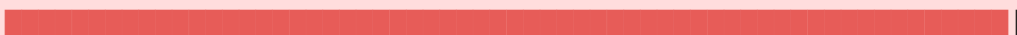
100%|  | 157/157
[00:00<00:00, 663.23it/s]
Test loss: 1.7108
Test accuracy: 0.7616


Epoka: 59/100

100%|  | 938/938
[00:03<00:00, 238.73it/s]
Train loss: 1.7801
Train accuracy: 0.6948


100%|  | 157/157
[00:00<00:00, 674.06it/s]
Test loss: 1.7094
Test accuracy: 0.7616


Epoka: 60/100

100%|  | 938/938
[00:03<00:00, 239.13it/s]
Train loss: 1.7777
Train accuracy: 0.6959

100%|  | 157/157
[00:00<00:00, 675.39it/s]
Test loss: 1.7076
Test accuracy: 0.7636

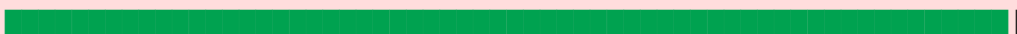
Epoka: 61/100

100%|  | 938/938
[00:03<00:00, 239.61it/s]
Train loss: 1.7759
Train accuracy: 0.6976

100%|  | 157/157
[00:00<00:00, 678.15it/s]
Test loss: 1.7060
Test accuracy: 0.7648

Epoka: 62/100


100%|  | 938/938
[00:03<00:00, 235.24it/s]
Train loss: 1.7717
Train accuracy: 0.7012

100%|  | 157/157
[00:00<00:00, 593.69it/s]
Test loss: 1.7055
Test accuracy: 0.7652


Epoka: 63/100


100%|  | 938/938
[00:03<00:00, 239.62it/s]

Train loss: 1.7699
Train accuracy: 0.7037


100%|  | 157/157
[00:00<00:00, 673.99it/s]
Test loss: 1.7042
Test accuracy: 0.7660

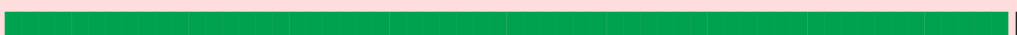
Epoka: 64/100

100%|  | 938/938
[00:03<00:00, 239.55it/s]
Train loss: 1.7660
Train accuracy: 0.7076

100%|  | 157/157
[00:00<00:00, 676.79it/s]
Test loss: 1.7036
Test accuracy: 0.7658


Epoka: 65/100

100%|  | 938/938
[00:03<00:00, 236.60it/s]
Train loss: 1.7656
Train accuracy: 0.7071

100%|  | 157/157
[00:00<00:00, 663.19it/s]
Test loss: 1.7011
Test accuracy: 0.7684

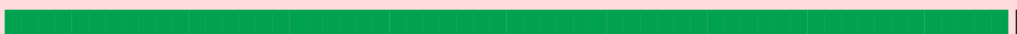
Epoka: 66/100

100%|  | 938/938
[00:03<00:00, 240.33it/s]
Train loss: 1.7634
Train accuracy: 0.7099

100%|  | 157/157
[00:00<00:00, 644.84it/s]
Test loss: 1.7019
Test accuracy: 0.7674

Epoka: 67/100


100%|  | 938/938
[00:03<00:00, 237.69it/s]
Train loss: 1.7615
Train accuracy: 0.7101

100%|  | 157/157
[00:00<00:00, 675.59it/s]
Test loss: 1.6993
Test accuracy: 0.7701


Epoka: 68/100


100%|  | 938/938
[00:03<00:00, 237.88it/s]

Train loss: 1.7595
Train accuracy: 0.7129


100%|  | 157/157
[00:00<00:00, 673.30it/s]
Test loss: 1.6986
Test accuracy: 0.7705

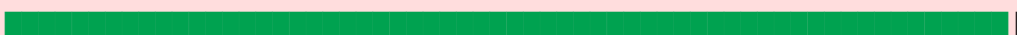
Epoka: 69/100

100%|  | 938/938
[00:03<00:00, 240.28it/s]
Train loss: 1.7563
Train accuracy: 0.7152

100%|  | 157/157
[00:00<00:00, 666.92it/s]
Test loss: 1.6978
Test accuracy: 0.7714


Epoka: 70/100

100%|  | 938/938
[00:03<00:00, 238.14it/s]
Train loss: 1.7568
Train accuracy: 0.7140


100%|  | 157/157
[00:00<00:00, 678.51it/s]
Test loss: 1.6965
Test accuracy: 0.7731

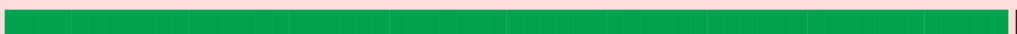
Epoka: 71/100

100%|  | 938/938
[00:03<00:00, 239.67it/s]
Train loss: 1.7542
Train accuracy: 0.7171

100%|  | 157/157
[00:00<00:00, 675.45it/s]
Test loss: 1.6955
Test accuracy: 0.7723

Epoka: 72/100


100%|  | 938/938
[00:03<00:00, 239.72it/s]
Train loss: 1.7524
Train accuracy: 0.7183

100%|  | 157/157
[00:00<00:00, 679.26it/s]
Test loss: 1.6945
Test accuracy: 0.7740


Epoka: 73/100

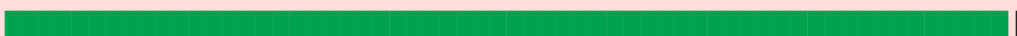
100%|  | 938/938
[00:03<00:00, 238.92it/s]

Train loss: 1.7503
Train accuracy: 0.7217

100%|  | 157/157
[00:00<00:00, 689.95it/s]
Test loss: 1.6936
Test accuracy: 0.7738

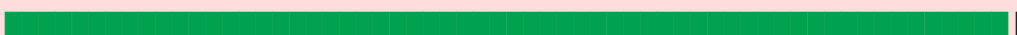
Epoka: 74/100

100%|  | 938/938
[00:03<00:00, 238.95it/s]
Train loss: 1.7499
Train accuracy: 0.7205

100%|  | 157/157
[00:00<00:00, 682.48it/s]
Test loss: 1.6936
Test accuracy: 0.7744

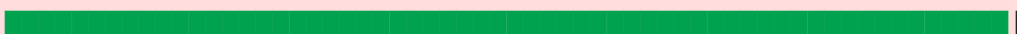
Epoka: 75/100

100%|  | 938/938
[00:03<00:00, 235.66it/s]
Train loss: 1.7496
Train accuracy: 0.7210

100%|  | 157/157
[00:00<00:00, 602.87it/s]
Test loss: 1.6933
Test accuracy: 0.7756

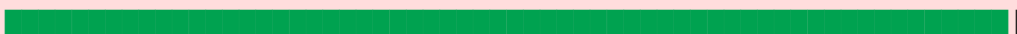
Epoka: 76/100

100%|  | 938/938
[00:03<00:00, 239.24it/s]
Train loss: 1.7469
Train accuracy: 0.7240

100%|  | 157/157
[00:00<00:00, 675.33it/s]
Test loss: 1.6929
Test accuracy: 0.7742

Epoka: 77/100


100%|  | 938/938
[00:03<00:00, 239.56it/s]
Train loss: 1.7459
Train accuracy: 0.7252

100%|  | 157/157
[00:00<00:00, 617.45it/s]
Test loss: 1.6906
Test accuracy: 0.7773


Epoka: 78/100

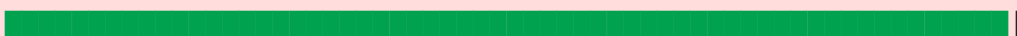
100%|  | 938/938
[00:03<00:00, 237.46it/s]

Train loss: 1.7430
Train accuracy: 0.7267

100%|  | 157/157
[00:00<00:00, 663.36it/s]
Test loss: 1.6900
Test accuracy: 0.7764

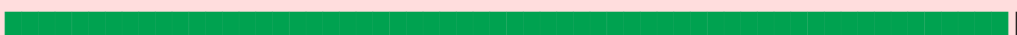
Epoka: 79/100

100%|  | 938/938
[00:03<00:00, 239.28it/s]
Train loss: 1.7439
Train accuracy: 0.7259

100%|  | 157/157
[00:00<00:00, 669.20it/s]
Test loss: 1.6897
Test accuracy: 0.7773

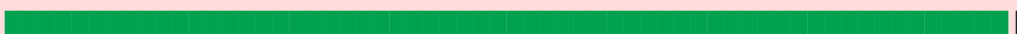
Epoka: 80/100

100%|  | 938/938
[00:03<00:00, 240.12it/s]
Train loss: 1.7415
Train accuracy: 0.7280

100%|  | 157/157
[00:00<00:00, 681.43it/s]
Test loss: 1.6891
Test accuracy: 0.7781

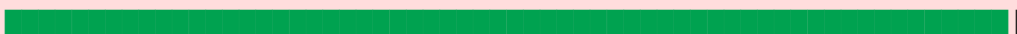
Epoka: 81/100

100%|  | 938/938
[00:03<00:00, 237.58it/s]
Train loss: 1.7412
Train accuracy: 0.7287

100%|  | 157/157
[00:00<00:00, 674.95it/s]
Test loss: 1.6889
Test accuracy: 0.7778

Epoka: 82/100


100%|  | 938/938
[00:03<00:00, 239.24it/s]
Train loss: 1.7393
Train accuracy: 0.7301

100%|  | 157/157
[00:00<00:00, 678.86it/s]
Test loss: 1.6886
Test accuracy: 0.7785


Epoka: 83/100

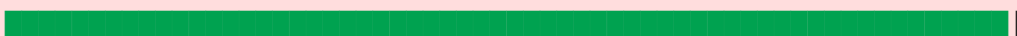
100%|  | 938/938
[00:03<00:00, 239.88it/s]

Train loss: 1.7385
Train accuracy: 0.7310

100%|  | 157/157
[00:00<00:00, 602.86it/s]
Test loss: 1.6877
Test accuracy: 0.7787

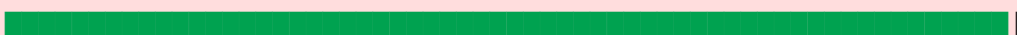
Epoka: 84/100

100%|  | 938/938
[00:03<00:00, 238.88it/s]
Train loss: 1.7376
Train accuracy: 0.7315

100%|  | 157/157
[00:00<00:00, 677.97it/s]
Test loss: 1.6876
Test accuracy: 0.7792


Epoka: 85/100

100%|  | 938/938
[00:03<00:00, 239.57it/s]
Train loss: 1.7356
Train accuracy: 0.7337

100%|  | 157/157
[00:00<00:00, 680.47it/s]
Test loss: 1.6866
Test accuracy: 0.7796

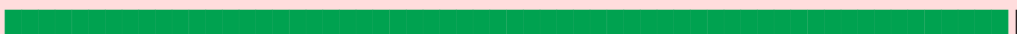
Epoka: 86/100

100%|  | 938/938
[00:03<00:00, 237.68it/s]
Train loss: 1.7353
Train accuracy: 0.7339

100%|  | 157/157
[00:00<00:00, 664.36it/s]
Test loss: 1.6857
Test accuracy: 0.7802

Epoka: 87/100


100%|  | 938/938
[00:03<00:00, 240.11it/s]
Train loss: 1.7336
Train accuracy: 0.7353

100%|  | 157/157
[00:00<00:00, 674.27it/s]
Test loss: 1.6854
Test accuracy: 0.7809


Epoka: 88/100


100%|  | 938/938
[00:03<00:00, 240.54it/s]

Train loss: 1.7338
Train accuracy: 0.7346

100%|  | 157/157
[00:00<00:00, 649.29it/s]
Test loss: 1.6842
Test accuracy: 0.7814

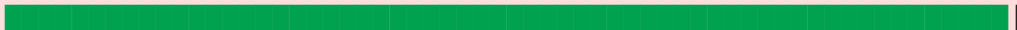
Epoka: 89/100

100%|  | 938/938
[00:03<00:00, 238.05it/s]
Train loss: 1.7324
Train accuracy: 0.7358

100%|  | 157/157
[00:00<00:00, 669.34it/s]
Test loss: 1.6842
Test accuracy: 0.7819

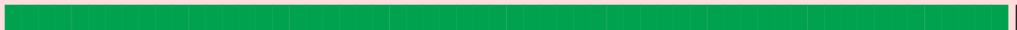
Epoka: 90/100

100%|  | 938/938
[00:03<00:00, 237.74it/s]
Train loss: 1.7309
Train accuracy: 0.7379


100%|  | 157/157
[00:00<00:00, 659.02it/s]
Test loss: 1.6843
Test accuracy: 0.7813


Epoka: 91/100

100%|  | 938/938
[00:03<00:00, 237.96it/s]
Train loss: 1.7311
Train accuracy: 0.7370

100%|  | 157/157
[00:00<00:00, 680.71it/s]
Test loss: 1.6836
Test accuracy: 0.7820

Epoka: 92/100


100%|  | 938/938
[00:03<00:00, 238.76it/s]
Train loss: 1.7287
Train accuracy: 0.7398

100%|  | 157/157
[00:00<00:00, 661.63it/s]
Test loss: 1.6824
Test accuracy: 0.7822

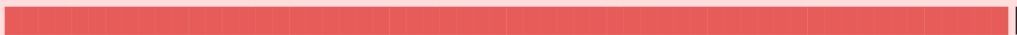
Epoka: 93/100


100%|  | 938/938
[00:03<00:00, 238.81it/s]

Train loss: 1.7278
Train accuracy: 0.7407

100%|  | 157/157
[00:00<00:00, 673.09it/s]
Test loss: 1.6827
Test accuracy: 0.7825


Epoka: 94/100

100%|  | 938/938
[00:03<00:00, 238.77it/s]
Train loss: 1.7273
Train accuracy: 0.7410


100%|  | 157/157
[00:00<00:00, 676.83it/s]
Test loss: 1.6820
Test accuracy: 0.7825


Epoka: 95/100

100%|  | 938/938
[00:03<00:00, 239.07it/s]
Train loss: 1.7275
Train accuracy: 0.7405

100%|  | 157/157
[00:00<00:00, 671.91it/s]
Test loss: 1.6824
Test accuracy: 0.7823

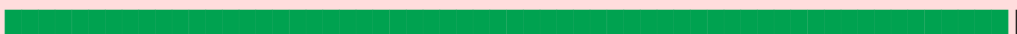
Epoka: 96/100

100%|  | 938/938
[00:03<00:00, 238.67it/s]
Train loss: 1.7259
Train accuracy: 0.7422

100%|  | 157/157
[00:00<00:00, 606.60it/s]
Test loss: 1.6822
Test accuracy: 0.7819

Epoka: 97/100

100%|  | 938/938
[00:04<00:00, 223.24it/s]
Train loss: 1.7256
Train accuracy: 0.7417

100%|  | 157/157
[00:00<00:00, 581.40it/s]
Test loss: 1.6805
Test accuracy: 0.7837

Epoka: 98/100

100%|  | 938/938
[00:03<00:00, 238.74it/s]

```
100%|██████████████████████████████████████████████████████████████████████████| 157/157  
[00:00<00:00, 669.25it/s]  
Test loss: 1.6809  
Test accuracy: 0.7830
```

```
100%|██████████████████████████████████████████████████████████████████████████| 938/938  
[00:03<00:00, 238.04it/s]  
Train loss: 1.7244  
Train accuracy: 0.7435
```

```
100%|██████████████████████████████████████████████████████████████████████████| 157/157  
[00:00<00:00, 670.58it/s]  
Test loss: 1.6811  
Test accuracy: 0.7826
```

```
100%|██████████████████████████████████████████████████████████████████████████| 938/938  
[00:03<00:00, 238.64it/s]  
Train loss: 1.7234  
Train accuracy: 0.7443
```

```
100%|██████████████████████████████████████████████████████████████████████████| 157/157  
[00:00<00:00, 675.67it/s]  
Test loss: 1.6821  
Test accuracy: 0.7815
```

```
In [27]: headers = history.keys()
         rows = zip(*history.values())
         print(tabulate(rows, headers=headers, tablefmt='pretty'))
```

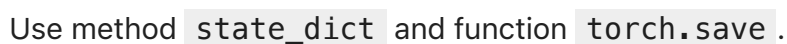
	loss_train	acc_train	loss_test	acc_test
2.302253891372681	0.11513333333333334	2.3015244480133057	0.1905	
2.3015840198516844	0.1659	2.300276218032837	0.2611	
2.300844240442912	0.16923333333333335	2.2987437141418456	0.217	
2.29984623743693	0.13713333333333333	2.2966318099975584	0.137	
2.2984486667633055	0.10833333333333334	2.2932555164337156	0.1096	
2.2958373302459716	0.10065	2.2861803852081297	0.0988	
2.290309753545125	0.09903333333333333	2.2713675563812257	0.0989	
2.282408146540324	0.11401666666666667	2.2551241802215576	0.1885	

2.274461056772868	0.1557	2.2402670627593992	0.2304
2.2661020184834797	0.18236666666666668	2.226756703186035	0.2549
2.2579274979909263	0.2041	2.2117916984558104	0.27
2.246191771952311	0.22096666666666667	2.188144364929199	0.2797
2.2321220471700034	0.22765	2.1603829486846924	0.2811
2.214310261408488	0.23635	2.1322396579742433	0.3105
2.195043679936727	0.26733333333333333	2.1030152755737306	0.3761
2.171557197189331	0.30548333333333333	2.0693416980743407	0.4537
2.1495066078186036	0.34176666666666666	2.032821955490112	0.5003
2.125733141454061	0.36795	1.9980935123443604	0.5231
2.107855320549011	0.38155	1.972504916191101	0.5324
2.089513953145345	0.39476666666666665	1.9505500785827636	0.5425
2.077560060755412	0.40123333333333333	1.9338432062149047	0.5478
2.0644258569081626	0.41318333333333335	1.9105321641921997	0.5792
2.051775353240967	0.42678333333333335	1.886344721031189	0.6171
2.0395387331644694	0.43783333333333335	1.8686718469619752	0.6237
2.02718345896403	0.45243333333333335	1.845398878288269	0.6784
2.012774441019694	0.47103333333333336	1.831066000556946	0.6865
2.0012552225748697	0.48278333333333334	1.820034048652649	0.6917
1.9900975998560588	0.49421666666666667	1.8101621887207031	0.698
1.9806062050501505	0.50268333333333334	1.8058312572479247	0.6989
1.9696693993886312	0.51341666666666666	1.7973448667526246	0.703
1.9600650800704955	0.52151666666666666	1.7917962175369262	0.7092
1.9525055934270223	0.52813333333333333	1.784802610397339	0.7122
1.942646544011434	0.53835	1.7837140434265137	0.7116
1.9323185910542806	0.54863333333333333	1.7763620609283448	0.7173

1.924548531850179	0.55535	1.7726523147583009	0.7222
1.9151218688964844	0.5648333333333333	1.7691554153442384	0.7247
1.9070255125045776	0.5748833333333333	1.7650784734725953	0.7288
1.8973816565831503	0.5829166666666666	1.7597935619354248	0.7316
1.8892160735448202	0.59315	1.7547349527359009	0.7368
1.8795310118993123	0.6014833333333334	1.7530566568374635	0.7368
1.8691610275268555	0.6125666666666667	1.7487932538986206	0.7379
1.862936517270406	0.6184666666666667	1.7470100233078003	0.7393
1.8532951913833617	0.6287666666666667	1.74525532207489	0.7386
1.8465696384429933	0.6353166666666666	1.741640721130371	0.7417
1.8415825243631998	0.6397166666666667	1.7383205612182617	0.7448
1.8340559964497885	0.6462166666666667	1.73600588722229	0.7467
1.8300407899220785	0.6486833333333333	1.7332159881591798	0.7476
1.8217039003372193	0.6572833333333333	1.7296825006484986	0.7495
1.8161588084538778	0.6636	1.726840637397766	0.751
1.8118267523447673	0.66705	1.7245335115432738	0.7527
1.807637751897176	0.6697166666666666	1.7231371269226075	0.754
1.8018449128468832	0.6755833333333333	1.721068071937561	0.7543
1.7996356215159097	0.6776666666666666	1.7185642642974854	0.7573
1.7941316501617433	0.6825333333333333	1.718650922203064	0.7547
1.7921761681874593	0.6845833333333333	1.7142424886703491	0.76
1.7888977373758952	0.6861333333333334	1.7142597557067871	0.7592
1.7845726915359497	0.69075	1.7111425048828126	0.762
1.782200001525879	0.6927166666666666	1.7108277946472168	0.7616
1.780063365618388	0.6948333333333333	1.7093623094558716	0.7616
1.7777117013295491	0.6958833333333333	1.7076313745498657	0.7636
1.775894584274292	0.6976333333333333	1.7059596210479737	0.7648

1.7717196739196777	0.7012	1.70550514087677	0.7652
1.7699199912389119	0.7036666666666667	1.7042118900299073	0.766
1.7660190184911093	0.7076166666666667	1.70364856300354	0.7658
1.7655882168451944	0.7071166666666666	1.7010777143478393	0.7684
1.7634079926172892	0.7098833333333333	1.7019110918045044	0.7674
1.7615486763636272	0.7100666666666666	1.6993078481674195	0.7701
1.7595406878789266	0.7128666666666666	1.6986120492935182	0.7705
1.7562510000864664	0.7151833333333333	1.6978486322402955	0.7714
1.7568376184463501	0.7140166666666666	1.6964753234863281	0.7731
1.7542193012873333	0.7171	1.6955408672332763	0.7723
1.752369516436259	0.7182666666666667	1.694457053565979	0.774
1.7502541810353598	0.7217	1.6935770433425903	0.7738
1.7499465909322103	0.72045	1.693608193206787	0.7744
1.7495777379989623	0.721	1.693255319404602	0.7756
1.7469037206013998	0.7239666666666666	1.6928973207473754	0.7742
1.7458859028498332	0.7252333333333333	1.6906259256362914	0.7773
1.7430447369893391	0.7266833333333333	1.6899588235855103	0.7764
1.7438965025583903	0.7259166666666667	1.6897081899642945	0.7773
1.741455255762736	0.728	1.6891486711502075	0.7781
1.7411678303400675	0.7287	1.6889302242279052	0.7778
1.739330744934082	0.7301333333333333	1.6886378028869629	0.7785
1.7384634753545125	0.7310166666666666	1.6876980792999268	0.7787
1.7376301279703776	0.7315333333333334	1.6876182680130005	0.7792
1.7355824155171713	0.7337	1.6866435876846313	0.7796
1.7353393600463867	0.7339333333333333	1.6856848493576049	0.7802
1.7336210536956786	0.7353166666666666	1.6854048727035522	0.7809

```
epochs = range(1, len(history["loss_train"]) + 1)
draw_loss_test(epochs, history)
```




```
In [29]: model_path = './models/net.pth'
optimizer_path = './models/optimizer.pth'

torch.save(net.state_dict(), model_path)
torch.save(optimizer.state_dict(), optimizer_path)
```

```
In [30]: model_path2 = './models/net2.pth'
optimizer_path2 = './models/optimizer2.pth'

torch.save(net2.state_dict(), model_path2)
torch.save(optimizer2.state_dict(), optimizer_path2)
```

```
In [31]: checkpoint_path = './checkpoints/checkpoint.pth'

torch.save(
    {
        "model_state_dict": net.state_dict(),
        "optimizer_state_dict": optimizer.state_dict(),
    },
    checkpoint_path,
)
```

```
In [32]: checkpoint_path2 = './checkpoints/checkpoint2.pth'

torch.save(
    {
        "model_state_dict": net2.state_dict(),
        "optimizer_state_dict": optimizer2.state_dict(),
    },
    checkpoint_path2,
)
```

13. Create new network with the same architecture and initialize it with saved weights. Compare evaluations for both networks.

Use `torch.load` and `load_state_dict`.

```
In [33]: net3 = MLP(input_shape, output_size)
optimizer3 = torch.optim.SGD(net2.parameters(), lr=0.01)

net3.load_state_dict(torch.load(model_path))
optimizer3.load_state_dict(torch.load(optimizer_path))
```

```
/var/folders/0b/brzkvl1j0tn9xzynh5pr39sw0000gn/T/ipykernel_17512/140974620
3.py:4: FutureWarning: You are using `torch.load` with `weights_only=False`
(the current default value), which uses the default pickle module implic
itly. It is possible to construct malicious pickle data which will execute
arbitrary code during unpickling (See https://github.com/pytorch/pytorch/b
lob/main/SECURITY.md#untrusted-models for more details). In a future relea
se, the default value for `weights_only` will be flipped to `True`. This l
imits the functions that could be executed during unpickling. Arbitrary ob
jects will no longer be allowed to be loaded via this mode unless they are
explicitly allowlisted by the user via `torch.serialization.add_safe_globa
ls`. We recommend you start setting `weights_only=True` for any use case w
here you don't have full control of the loaded file. Please open an issue
on GitHub for any issues related to this experimental feature.
```

```
net3.load_state_dict(torch.load(model_path))
```

```
/var/folders/0b/brzkvl1j0tn9xzynh5pr39sw0000gn/T/ipykernel_17512/140974620
3.py:5: FutureWarning: You are using `torch.load` with `weights_only=False`
(the current default value), which uses the default pickle module implic
itly. It is possible to construct malicious pickle data which will execute
arbitrary code during unpickling (See https://github.com/pytorch/pytorch/b
lob/main/SECURITY.md#untrusted-models for more details). In a future relea
se, the default value for `weights_only` will be flipped to `True`. This l
imits the functions that could be executed during unpickling. Arbitrary ob
jects will no longer be allowed to be loaded via this mode unless they are
explicitly allowlisted by the user via `torch.serialization.add_safe_globa
ls`. We recommend you start setting `weights_only=True` for any use case w
here you don't have full control of the loaded file. Please open an issue
on GitHub for any issues related to this experimental feature.
```

```
optimizer3.load_state_dict(torch.load(optimizer_path))
```

```
In [34]: checkpoint = torch.load("./checkpoints/checkpoint.pth")
net4 = MLP(input_shape, output_size)
optimizer4 = torch.optim.SGD(net3.parameters(), lr=0.01)

net4.load_state_dict(checkpoint["model_state_dict"])
optimizer4.load_state_dict(checkpoint["optimizer_state_dict"])

net4.to(device)
```

```
/var/folders/0b/brzkvl1j0tn9xzynh5pr39sw0000gn/T/ipykernel_17512/47768843.
py:1: FutureWarning: You are using `torch.load` with `weights_only=False`
(the current default value), which uses the default pickle module implic
itly. It is possible to construct malicious pickle data which will execute a
rbitrary code during unpickling (See https://github.com/pytorch/pytorch/bl
ob/main/SECURITY.md#untrusted-models for more details). In a future relea
se, the default value for `weights_only` will be flipped to `True`. This li
mits the functions that could be executed during unpickling. Arbitrary obj
ects will no longer be allowed to be loaded via this mode unless they are
explicitly allowlisted by the user via `torch.serialization.add_safe_globa
ls`. We recommend you start setting `weights_only=True` for any use case w
here you don't have full control of the loaded file. Please open an issue
on GitHub for any issues related to this experimental feature.
```

```
checkpoint = torch.load("./checkpoints/checkpoint.pth")
```

```
Out[34]: MLP(
  (model): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=784, out_features=512, bias=True)
    (2): ReLU()
    (3): Linear(in_features=512, out_features=512, bias=True)
    (4): ReLU()
    (5): Linear(in_features=512, out_features=10, bias=True)
    (6): Softmax(dim=1)
  )
)
```

14. EXTENSION EXERCISE

Define your own model and train it.

Try to achieve better results.

You can use different parameters, layers e.g.:

- conv2d
- maxpooling2d
- batch norm 2d
- and more...

Save weights to a file.

```
In [35]: import torch
import torch.nn as nn
import torch.optim as optim

class FastCNNModel(nn.Module):

    def __init__(self, input_shape, output_size) -> None:
        super(FastCNNModel, self).__init__()
        self.model = nn.Sequential(
            nn.Conv2d(
                in_channels=input_shape[0],
                out_channels=16,
                kernel_size=3,
                padding=1,
            ),
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2),
            nn.Conv2d(
                in_channels=16, out_channels=32, kernel_size=3, padding=1
            ),
            nn.BatchNorm2d(32),
            nn.ReLU(),
```

```
In [38]: my_net = FastCNNModel(input_shape, output_size)
my_net = my_net.to(device)

my_loss_fn = nn.CrossEntropyLoss()
my_optimizer = optim.SGD(my_net.parameters(), lr=0.01, momentum=0.9)
```

Epoch: 1/100

```
100%|████████████████████| 157/157 [00:00<00:00, 169.21it/s, acc
uracy=0.9783, loss=0.0637]
```


Epoch: 2/100

```
100%|████████████████████| 157/157 [00:01<00:00, 134.75it/s, acc
uracy=0.9860, loss=0.0429]
```

Epoch: 3/100


```
100%|████████████████████████████████████████| 157/157 [00:01<00:00, 140.04it/s, acc
uracy=0.9883, loss=0.0335]
```

Epoch: 4/100

100%|  | 938/938 [00:21<00:00, 43.52it/s, accuracy=0.9693, loss=0.0988]


100%|  | 157/157 [00:01<00:00, 147.58it/s, accuracy=0.9892, loss=0.0292]

Epoch: 5/100

100%|  | 938/938 [00:22<00:00, 42.43it/s, accuracy=0.9716, loss=0.0894]

100%|  | 157/157 [00:01<00:00, 140.32it/s, accuracy=0.9881, loss=0.0357]

Epoch: 6/100

100%|  | 938/938 [00:20<00:00, 46.21it/s, accuracy=0.9748, loss=0.0803]


100%|  | 157/157 [00:01<00:00, 147.13it/s, accuracy=0.9887, loss=0.0346]

Epoch: 7/100

100%|  | 938/938 [00:20<00:00, 45.93it/s, accuracy=0.9766, loss=0.0742]


100%|  | 157/157 [00:00<00:00, 157.42it/s, accuracy=0.9888, loss=0.0337]

Epoch: 8/100

100%|  | 938/938 [00:20<00:00, 45.29it/s, accuracy=0.9785, loss=0.0697]

100%|  | 157/157 [00:01<00:00, 145.98it/s, accuracy=0.9925, loss=0.0246]

Epoch: 9/100

100%|  | 938/938 [00:21<00:00, 43.49it/s, accuracy=0.9788, loss=0.0681]


100%|  | 157/157 [00:00<00:00, 181.97it/s, accuracy=0.9913, loss=0.0263]

Epoch: 10/100

100%|  | 938/938 [00:20<00:00, 46.66it/s, accuracy=0.9799, loss=0.0651]

100%|  | 157/157 [00:00<00:00, 191.42it/s, accuracy=0.9921, loss=0.0241]


Epoch: 11/100

100%|  | 938/938 [00:18<00:00, 50.54it/s, accuracy=0.9803, loss=0.0627]

100%|  | 157/157 [00:00<00:00, 185.16it/s, accuracy=0.9924, loss=0.0248]

Training accuracy of 0.9803 achieved, stopping training.

```
In [40]: my_epochs = range(1, len(my_history["loss_train"]) + 1)
         draw_loss_test(my_epochs, my_history)
```



```
In [41]: class SimpleFastCNNModel(nn.Module):

    def __init__(self, input_shape, output_size) -> None:
        super(SimpleFastCNNModel, self).__init__()
        self.model = nn.Sequential(
            nn.Conv2d(
                in_channels=input_shape[0],
                out_channels=8,
                kernel_size=3,
                padding=1,
            ),
            nn.BatchNorm2d(8),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2),
            nn.Conv2d(
                in_channels=8, out_channels=16, kernel_size=3, padding=1
            ),
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2),
            nn.AdaptiveAvgPool2d(1),
            nn.Flatten(),
            nn.Linear(16, output_size),
        )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.model(x)
```

```
In [45]: my_net2 = SimpleFastCNNModel(input_shape, output_size)
my_net2 = my_net2.to(device)

my_loss_fn2 = nn.CrossEntropyLoss()
my_optimizer2 = optim.SGD(my_net2.parameters(), lr=0.01, momentum=0.9)
```

```
In [46]: my_net2, my_history2 = test_or_train(
    my_net2,
```

Epoch: 1/100

```
100%|████████████████████| 157/157 [00:00<00:00, 291.11it/s, acc
uracy=0.6191, loss=1.2049]
```

```
100%|██████████| 938/938 [00:11<00:00, 78.98it/s, acc  
uracy=0.6599, loss=1.0627]
```

```
100%|████████████████████| 157/157 [00:00<00:00, 286.25it/s, acc
uracy=0.7627, loss=0.7199]
```

```
100%|██████████| 938/938 [00:11<00:00, 79.12it/s, acc  
uracy=0.7495, loss=0.7874]
```

```
100%|████████████████████████████████████████| 157/157 [00:00<00:00, 284.97it/s, acc
uracy=0.7657, loss=0.6997]
```

```
100%|████████████████████████████████████████| 938/938 [00:11<00:00, 79.11it/s, acc
uracy=0.7990, loss=0.6453]
```

```
100%|████████████████████████████████████████| 157/157 [00:00<00:00, 268.17it/s, acc
uracy=0.8615, loss=0.4619]
```

```
100%|██████████| 938/938 [00:12<00:00, 76.77it/s, acc  
uracy=0.8203, loss=0.5798]
```

```
100%|████████████████████████████████████████| 157/157 [00:00<00:00, 274.58it/s, acc
accuracy=0.8783, loss=0.4097]
```



```
100%|██████████| 938/938 [00:11<00:00, 78.68it/s, acc  
uracy=0.8342, loss=0.5287]
```



```
100%|████████████████████████████████████████| 157/157 [00:00<00:00, 277.01it/s, accuracy=0.9102, loss=0.3242]
```



```
100%|████████████████████████████████████████| 938/938 [00:11<00:00, 79.33it/s, acc
uracy=0.8514, loss=0.4839]
```



```
100%|████████████████████████████████████████| 157/157 [00:00<00:00, 285.39it/s, acc
uracy=0.8951, loss=0.3536]
```



Epoch: 8/100



100%|  | 938/938 [00:11<00:00, 79.34it/s, accuracy=0.8563, loss=0.4606]
100%|  | 157/157 [00:00<00:00, 287.09it/s, accuracy=0.8962, loss=0.3302]
Epoch: 9/100



100%|  | 938/938 [00:11<00:00, 79.44it/s, accuracy=0.8641, loss=0.4383]
100%|  | 157/157 [00:00<00:00, 283.92it/s, accuracy=0.9130, loss=0.2996]
Epoch: 10/100


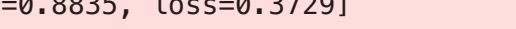
100%|  | 938/938 [00:12<00:00, 77.43it/s, accuracy=0.8658, loss=0.4287]
100%|  | 157/157 [00:00<00:00, 285.38it/s, accuracy=0.9286, loss=0.2466]
Epoch: 11/100



100%|  | 938/938 [00:11<00:00, 79.31it/s, accuracy=0.8717, loss=0.4103]
100%|  | 157/157 [00:00<00:00, 290.86it/s, accuracy=0.9125, loss=0.2946]
Epoch: 12/100


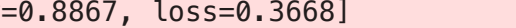
100%|  | 938/938 [00:11<00:00, 79.35it/s, accuracy=0.8781, loss=0.3975]
100%|  | 157/157 [00:00<00:00, 286.46it/s, accuracy=0.9172, loss=0.2728]
Epoch: 13/100



100%|  | 938/938 [00:11<00:00, 79.19it/s, accuracy=0.8779, loss=0.3936]
100%|  | 157/157 [00:00<00:00, 289.21it/s, accuracy=0.9283, loss=0.2346]
Epoch: 14/100



100%|  | 938/938 [00:11<00:00, 79.12it/s, accuracy=0.8764, loss=0.3931]
100%|  | 157/157 [00:00<00:00, 287.19it/s, accuracy=0.9288, loss=0.2365]
Epoch: 15/100



100%|  | 938/938 [00:11<00:00, 79.14it/s, accuracy=0.8835, loss=0.3729]
100%|  | 157/157 [00:00<00:00, 264.25it/s, accuracy=0.9336, loss=0.2269]
Epoch: 16/100



100%|  | 938/938 [00:11<00:00, 79.29it/s, accuracy=0.8847, loss=0.3671]
100%|  | 157/157 [00:00<00:00, 265.41it/s, accuracy=0.9332, loss=0.2297]
Epoch: 17/100



100%|  | 938/938 [00:11<00:00, 78.92it/s, accuracy=0.8867, loss=0.3668]
100%|  | 157/157 [00:00<00:00, 279.61it/s, accuracy=0.9403, loss=0.2047]
Epoch: 18/100



100%|  | 938/938 [00:11<00:00, 79.48it/s, accuracy=0.8895, loss=0.3560]
100%|  | 157/157 [00:00<00:00, 289.79it/s, accuracy=0.9318, loss=0.2331]
Epoch: 19/100



100%|  | 938/938 [00:11<00:00, 80.07it/s, accuracy=0.8928, loss=0.3499]
100%|  | 157/157 [00:00<00:00, 290.81it/s, accuracy=0.9371, loss=0.2140]
Epoch: 20/100


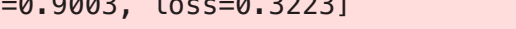
100%|  | 938/938 [00:11<00:00, 80.02it/s, accuracy=0.8926, loss=0.3443]
100%|  | 157/157 [00:00<00:00, 290.61it/s, accuracy=0.9316, loss=0.2264]
Epoch: 21/100



100%|  | 938/938 [00:11<00:00, 79.61it/s, accuracy=0.8928, loss=0.3455]
100%|  | 157/157 [00:00<00:00, 286.85it/s, accuracy=0.9344, loss=0.2143]
Epoch: 22/100


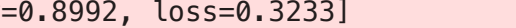
100%|  | 938/938 [00:11<00:00, 79.80it/s, accuracy=0.8948, loss=0.3392]
100%|  | 157/157 [00:00<00:00, 290.39it/s, accuracy=0.9379, loss=0.2173]
Epoch: 23/100



100%|  | 938/938 [00:11<00:00, 78.67it/s, accuracy=0.8957, loss=0.3343]
100%|  | 157/157 [00:00<00:00, 277.56it/s, accuracy=0.9457, loss=0.1908]
Epoch: 24/100



100%|  | 938/938 [00:11<00:00, 80.10it/s, accuracy=0.8983, loss=0.3302]
100%|  | 157/157 [00:00<00:00, 279.19it/s, accuracy=0.9342, loss=0.2181]
Epoch: 25/100



100%|  | 938/938 [00:11<00:00, 80.23it/s, accuracy=0.9003, loss=0.3223]
100%|  | 157/157 [00:00<00:00, 295.19it/s, accuracy=0.9448, loss=0.1879]
Epoch: 26/100



100%|  | 938/938 [00:39<00:00, 23.76it/s, accuracy=0.8994, loss=0.3249]
100%|  | 157/157 [00:00<00:00, 238.36it/s, accuracy=0.9498, loss=0.1875]
Epoch: 27/100



100%|  | 938/938 [00:12<00:00, 76.69it/s, accuracy=0.8992, loss=0.3233]
100%|  | 157/157 [00:00<00:00, 280.33it/s, accuracy=0.9408, loss=0.1947]
Epoch: 28/100



100%|  | 938/938 [00:11<00:00, 78.66it/s, accuracy=0.8994, loss=0.3240]
100%|  | 157/157 [00:00<00:00, 266.86it/s, accuracy=0.9455, loss=0.1863]
Epoch: 29/100



100%|  | 938/938 [00:12<00:00, 73.42it/s, accuracy=0.9013, loss=0.3180]
100%|  | 157/157 [00:00<00:00, 228.04it/s, accuracy=0.9473, loss=0.1765]
Epoch: 30/100


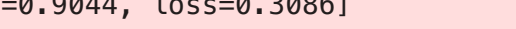
100%|  | 938/938 [00:11<00:00, 78.32it/s, accuracy=0.9009, loss=0.3175]
100%|  | 157/157 [00:00<00:00, 272.28it/s, accuracy=0.9401, loss=0.1966]
Epoch: 31/100



100%|  | 938/938 [00:11<00:00, 78.54it/s, accuracy=0.9046, loss=0.3128]
100%|  | 157/157 [00:00<00:00, 275.49it/s, accuracy=0.9487, loss=0.1775]
Epoch: 32/100


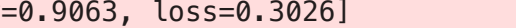
100%|  | 938/938 [00:11<00:00, 78.57it/s, accuracy=0.9039, loss=0.3097]
100%|  | 157/157 [00:00<00:00, 271.62it/s, accuracy=0.9476, loss=0.1753]
Epoch: 33/100



100%|  | 938/938 [00:11<00:00, 78.88it/s, accuracy=0.9028, loss=0.3137]
100%|  | 157/157 [00:00<00:00, 261.39it/s, accuracy=0.9442, loss=0.1904]
Epoch: 34/100



100%|  | 938/938 [00:12<00:00, 77.23it/s, accuracy=0.9043, loss=0.3091]
100%|  | 157/157 [00:00<00:00, 280.07it/s, accuracy=0.9366, loss=0.2017]
Epoch: 35/100



100%|  | 938/938 [00:11<00:00, 78.52it/s, accuracy=0.9044, loss=0.3086]
100%|  | 157/157 [00:00<00:00, 279.04it/s, accuracy=0.9470, loss=0.1807]
Epoch: 36/100



100%|  | 938/938 [00:11<00:00, 78.68it/s, accuracy=0.9043, loss=0.3066]
100%|  | 157/157 [00:00<00:00, 278.18it/s, accuracy=0.9395, loss=0.1869]
Epoch: 37/100



100%|  | 938/938 [00:11<00:00, 78.90it/s, accuracy=0.9063, loss=0.3026]
100%|  | 157/157 [00:00<00:00, 261.12it/s, accuracy=0.9395, loss=0.1987]
Epoch: 38/100



100%|  | 938/938 [00:11<00:00, 79.06it/s, accuracy=0.9060, loss=0.3025]
100%|  | 157/157 [00:00<00:00, 267.94it/s, accuracy=0.9444, loss=0.1877]
Epoch: 39/100



100%|  | 938/938 [00:11<00:00, 79.14it/s, accuracy=0.9078, loss=0.3014]
100%|  | 157/157 [00:00<00:00, 262.59it/s, accuracy=0.9442, loss=0.1802]
Epoch: 40/100


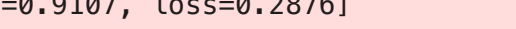
100%|  | 938/938 [00:11<00:00, 79.37it/s, accuracy=0.9058, loss=0.3007]
100%|  | 157/157 [00:00<00:00, 270.84it/s, accuracy=0.9515, loss=0.1634]
Epoch: 41/100


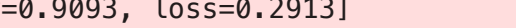
100%|  | 938/938 [00:11<00:00, 79.38it/s, accuracy=0.9082, loss=0.2974]
100%|  | 157/157 [00:00<00:00, 268.06it/s, accuracy=0.9503, loss=0.1633]
Epoch: 42/100


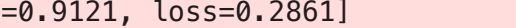
100%|  | 938/938 [00:11<00:00, 78.76it/s, accuracy=0.9080, loss=0.2975]
100%|  | 157/157 [00:00<00:00, 279.08it/s, accuracy=0.9545, loss=0.1551]
Epoch: 43/100



100%|  | 938/938 [00:11<00:00, 79.20it/s, accuracy=0.9086, loss=0.2965]
100%|  | 157/157 [00:00<00:00, 276.27it/s, accuracy=0.9538, loss=0.1607]
Epoch: 44/100



100%|  | 938/938 [00:11<00:00, 79.30it/s, accuracy=0.9090, loss=0.2916]
100%|  | 157/157 [00:00<00:00, 285.40it/s, accuracy=0.9495, loss=0.1726]
Epoch: 45/100



100%|  | 938/938 [00:11<00:00, 79.31it/s, accuracy=0.9107, loss=0.2876]
100%|  | 157/157 [00:00<00:00, 284.83it/s, accuracy=0.9443, loss=0.1797]
Epoch: 46/100



100%|  | 938/938 [00:11<00:00, 79.34it/s, accuracy=0.9093, loss=0.2913]
100%|  | 157/157 [00:00<00:00, 283.83it/s, accuracy=0.9464, loss=0.1765]
Epoch: 47/100



100%|  | 938/938 [00:12<00:00, 75.77it/s, accuracy=0.9121, loss=0.2861]
100%|  | 157/157 [00:00<00:00, 256.36it/s, accuracy=0.9498, loss=0.1664]
Epoch: 48/100



100%|  | 938/938 [00:11<00:00, 78.73it/s, accuracy=0.9101, loss=0.2876]
100%|  | 157/157 [00:00<00:00, 269.65it/s, accuracy=0.9518, loss=0.1637]
Epoch: 49/100



100%|  | 938/938 [00:11<00:00, 78.38it/s, accuracy=0.9133, loss=0.2798]
100%|  | 157/157 [00:00<00:00, 272.37it/s, accuracy=0.9447, loss=0.1825]
Epoch: 50/100



100%|  | 938/938 [00:11<00:00, 78.28it/s, accuracy=0.9123, loss=0.2822]
100%|  | 157/157 [00:00<00:00, 275.05it/s, accuracy=0.9492, loss=0.1633]
Epoch: 51/100



100%|  | 938/938 [00:11<00:00, 78.49it/s, accuracy=0.9132, loss=0.2766]
100%|  | 157/157 [00:00<00:00, 274.05it/s, accuracy=0.9496, loss=0.1666]
Epoch: 52/100



100%|  | 938/938 [00:12<00:00, 78.14it/s, accuracy=0.9142, loss=0.2793]
100%|  | 157/157 [00:00<00:00, 279.13it/s, accuracy=0.9498, loss=0.1651]
Epoch: 53/100



100%|  | 938/938 [00:11<00:00, 78.87it/s, accuracy=0.9133, loss=0.2802]
100%|  | 157/157 [00:00<00:00, 265.18it/s, accuracy=0.9549, loss=0.1522]
Epoch: 54/100



100%|  | 938/938 [00:11<00:00, 78.53it/s, accuracy=0.9164, loss=0.2704]
100%|  | 157/157 [00:00<00:00, 269.68it/s, accuracy=0.9501, loss=0.1656]
Epoch: 55/100



100%|  | 938/938 [00:11<00:00, 79.07it/s, accuracy=0.9140, loss=0.2793]
100%|  | 157/157 [00:00<00:00, 273.85it/s, accuracy=0.9552, loss=0.1503]
Epoch: 56/100



100%|  | 938/938 [00:11<00:00, 78.83it/s, accuracy=0.9141, loss=0.2798]
100%|  | 157/157 [00:00<00:00, 277.04it/s, accuracy=0.9497, loss=0.1705]
Epoch: 57/100



100%|  | 938/938 [00:11<00:00, 79.27it/s, accuracy=0.9131, loss=0.2776]
100%|  | 157/157 [00:00<00:00, 279.15it/s, accuracy=0.9550, loss=0.1505]
Epoch: 58/100



100%|  | 938/938 [00:11<00:00, 79.03it/s, accuracy=0.9156, loss=0.2740]
100%|  | 157/157 [00:00<00:00, 281.63it/s, accuracy=0.9516, loss=0.1548]
Epoch: 59/100



100%|  | 938/938 [00:11<00:00, 79.92it/s, accuracy=0.9160, loss=0.2710]
100%|  | 157/157 [00:00<00:00, 278.30it/s, accuracy=0.9539, loss=0.1544]
Epoch: 60/100



100%|  | 938/938 [00:11<00:00, 79.63it/s, accuracy=0.9153, loss=0.2707]
100%|  | 157/157 [00:00<00:00, 284.80it/s, accuracy=0.9538, loss=0.1531]
Epoch: 61/100



100%|  | 938/938 [00:11<00:00, 79.28it/s, accuracy=0.9155, loss=0.2725]
100%|  | 157/157 [00:00<00:00, 284.22it/s, accuracy=0.9558, loss=0.1523]
Epoch: 62/100



100%|  | 938/938 [00:11<00:00, 79.84it/s, accuracy=0.9145, loss=0.2767]
100%|  | 157/157 [00:00<00:00, 279.15it/s, accuracy=0.9558, loss=0.1474]
Epoch: 63/100



100%|  | 938/938 [00:11<00:00, 79.46it/s, accuracy=0.9161, loss=0.2694]
100%|  | 157/157 [00:00<00:00, 282.57it/s, accuracy=0.9504, loss=0.1637]
Epoch: 64/100



100%|  | 938/938 [00:12<00:00, 77.92it/s, accuracy=0.9164, loss=0.2710]
100%|  | 157/157 [00:00<00:00, 279.58it/s, accuracy=0.9531, loss=0.1521]
Epoch: 65/100



100%|  | 938/938 [00:11<00:00, 79.42it/s, accuracy=0.9165, loss=0.2671]
100%|  | 157/157 [00:00<00:00, 275.35it/s, accuracy=0.9518, loss=0.1575]
Epoch: 66/100



100%|  | 938/938 [00:11<00:00, 79.64it/s, accuracy=0.9188, loss=0.2664]
100%|  | 157/157 [00:00<00:00, 263.34it/s, accuracy=0.9586, loss=0.1440]
Epoch: 67/100



100%|  | 938/938 [00:12<00:00, 76.30it/s, accuracy=0.9179, loss=0.2653]
100%|  | 157/157 [00:00<00:00, 252.24it/s, accuracy=0.9519, loss=0.1524]
Epoch: 68/100



100%|  | 938/938 [00:11<00:00, 81.12it/s, accuracy=0.9172, loss=0.2642]
100%|  | 157/157 [00:00<00:00, 307.49it/s, accuracy=0.9568, loss=0.1405]
Epoch: 69/100



100%|  | 938/938 [00:11<00:00, 84.13it/s, accuracy=0.9197, loss=0.2608]
100%|  | 157/157 [00:00<00:00, 313.86it/s, accuracy=0.9509, loss=0.1619]
Epoch: 70/100


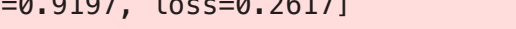
100%|  | 938/938 [00:11<00:00, 83.60it/s, accuracy=0.9193, loss=0.2637]
100%|  | 157/157 [00:00<00:00, 248.83it/s, accuracy=0.9503, loss=0.1581]
Epoch: 71/100


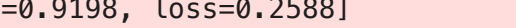
100%|  | 938/938 [00:11<00:00, 80.68it/s, accuracy=0.9195, loss=0.2611]
100%|  | 157/157 [00:00<00:00, 306.41it/s, accuracy=0.9576, loss=0.1441]
Epoch: 72/100


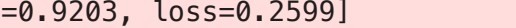
100%|  | 938/938 [00:11<00:00, 80.24it/s, accuracy=0.9196, loss=0.2637]
100%|  | 157/157 [00:00<00:00, 305.95it/s, accuracy=0.9552, loss=0.1471]
Epoch: 73/100



100%|  | 938/938 [00:11<00:00, 84.18it/s, accuracy=0.9211, loss=0.2557]
100%|  | 157/157 [00:00<00:00, 301.92it/s, accuracy=0.9552, loss=0.1468]
Epoch: 74/100



100%|  | 938/938 [00:10<00:00, 86.07it/s, accuracy=0.9188, loss=0.2623]
100%|  | 157/157 [00:00<00:00, 296.90it/s, accuracy=0.9590, loss=0.1391]
Epoch: 75/100



100%|  | 938/938 [00:11<00:00, 83.10it/s, accuracy=0.9197, loss=0.2617]
100%|  | 157/157 [00:00<00:00, 288.42it/s, accuracy=0.9587, loss=0.1356]
Epoch: 76/100



100%|  | 938/938 [00:11<00:00, 81.51it/s, accuracy=0.9198, loss=0.2588]
100%|  | 157/157 [00:00<00:00, 271.31it/s, accuracy=0.9547, loss=0.1466]
Epoch: 77/100



100%|  | 938/938 [00:12<00:00, 78.13it/s, accuracy=0.9203, loss=0.2599]
100%|  | 157/157 [00:00<00:00, 279.63it/s, accuracy=0.9579, loss=0.1417]
Epoch: 78/100



100%|  | 938/938 [00:11<00:00, 83.37it/s, accuracy=0.9197, loss=0.2574]
100%|  | 157/157 [00:00<00:00, 308.48it/s, accuracy=0.9563, loss=0.1487]
Epoch: 79/100



100%|  | 938/938 [00:11<00:00, 84.60it/s, accuracy=0.9193, loss=0.2587]
100%|  | 157/157 [00:00<00:00, 317.58it/s, accuracy=0.9508, loss=0.1600]
Epoch: 80/100



100%|  | 938/938 [00:11<00:00, 83.66it/s, accuracy=0.9193, loss=0.2603]
100%|  | 157/157 [00:00<00:00, 297.17it/s, accuracy=0.9536, loss=0.1526]
Epoch: 81/100



100%|  | 938/938 [00:11<00:00, 82.49it/s, accuracy=0.9226, loss=0.2498]
100%|  | 157/157 [00:00<00:00, 289.99it/s, accuracy=0.9576, loss=0.1396]
Epoch: 82/100



100%|  | 938/938 [00:11<00:00, 83.22it/s, accuracy=0.9196, loss=0.2582]
100%|  | 157/157 [00:00<00:00, 281.37it/s, accuracy=0.9523, loss=0.1561]
Epoch: 83/100



100%|  | 938/938 [00:11<00:00, 83.80it/s, accuracy=0.9210, loss=0.2559]
100%|  | 157/157 [00:00<00:00, 291.67it/s, accuracy=0.9577, loss=0.1429]
Epoch: 84/100



100%|  | 938/938 [00:11<00:00, 84.10it/s, accuracy=0.9206, loss=0.2565]
100%|  | 157/157 [00:00<00:00, 321.60it/s, accuracy=0.9603, loss=0.1345]
Epoch: 85/100



100%|  | 938/938 [00:11<00:00, 84.23it/s, accuracy=0.9227, loss=0.2518]
100%|  | 157/157 [00:00<00:00, 297.64it/s, accuracy=0.9544, loss=0.1461]
Epoch: 86/100



100%|  | 938/938 [00:11<00:00, 82.93it/s, accuracy=0.9218, loss=0.2553]
100%|  | 157/157 [00:00<00:00, 318.20it/s, accuracy=0.9548, loss=0.1464]
Epoch: 87/100



100%|  | 938/938 [00:11<00:00, 80.48it/s, accuracy=0.9214, loss=0.2518]
100%|  | 157/157 [00:00<00:00, 278.65it/s, accuracy=0.9566, loss=0.1394]
Epoch: 88/100



100%|  | 938/938 [00:11<00:00, 82.91it/s, accuracy=0.9228, loss=0.2494]
100%|  | 157/157 [00:00<00:00, 284.68it/s, accuracy=0.9606, loss=0.1341]
Epoch: 89/100



100%|  | 938/938 [00:11<00:00, 84.02it/s, accuracy=0.9235, loss=0.2495]
100%|  | 157/157 [00:00<00:00, 307.35it/s, accuracy=0.9598, loss=0.1373]
Epoch: 90/100


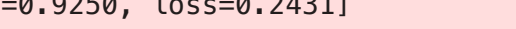
100%|  | 938/938 [00:11<00:00, 82.54it/s, accuracy=0.9232, loss=0.2499]
100%|  | 157/157 [00:00<00:00, 320.00it/s, accuracy=0.9595, loss=0.1363]
Epoch: 91/100


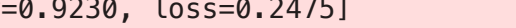
100%|  | 938/938 [00:11<00:00, 82.29it/s, accuracy=0.9232, loss=0.2505]
100%|  | 157/157 [00:00<00:00, 296.37it/s, accuracy=0.9585, loss=0.1366]
Epoch: 92/100


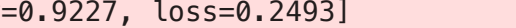
100%|  | 938/938 [00:11<00:00, 82.18it/s, accuracy=0.9236, loss=0.2459]
100%|  | 157/157 [00:00<00:00, 296.85it/s, accuracy=0.9590, loss=0.1372]
Epoch: 93/100

100%|  | 938/938 [00:11<00:00, 83.82it/s, accuracy=0.9247, loss=0.2484]
100%|  | 157/157 [00:00<00:00, 290.85it/s, accuracy=0.9588, loss=0.1363]
Epoch: 94/100

100%|  | 938/938 [00:11<00:00, 84.99it/s, accuracy=0.9234, loss=0.2461]
100%|  | 157/157 [00:00<00:00, 313.59it/s, accuracy=0.9602, loss=0.1311]
Epoch: 95/100

100%|  | 938/938 [00:11<00:00, 80.34it/s, accuracy=0.9250, loss=0.2431]
100%|  | 157/157 [00:00<00:00, 277.52it/s, accuracy=0.9570, loss=0.1383]
Epoch: 96/100

100%|  | 938/938 [00:12<00:00, 77.95it/s, accuracy=0.9230, loss=0.2475]
100%|  | 157/157 [00:00<00:00, 277.87it/s, accuracy=0.9596, loss=0.1384]
Epoch: 97/100

100%|  | 938/938 [00:11<00:00, 79.06it/s, accuracy=0.9227, loss=0.2493]
100%|  | 157/157 [00:00<00:00, 278.94it/s, accuracy=0.9585, loss=0.1373]
Epoch: 98/100


```
In [47]: my_epochs2 = range(1, len(my_history2["loss_train"]) + 1)
         draw_loss_test(my_epochs2, my_history2)
```

